

Figures/flowchart_part1.pdf

Figures/Taxonomy.jpg

Figure 2: **Taxonomy of LTU Attacker** . **Top:** Any LTU Attacker has available the Defender trainer $\mathcal{T}_D$, the trained Defender model $\mathcal{M}_D$, and attack data $\mathcal{D}_A$ including (almost) all the Defender data $\mathcal{D}_D$ and Reserved data $\mathcal{D}_R$ $\mathcal{D}_A = \mathcal{D}_D - \{\text{membership}(d)\} \cup \mathcal{D}_R - \{\text{membership}(r)\}$. But it may use only part of this available knowledge to conduct attacks. r and d are two labeled examples belonging $\mathcal{D}_R$ and $\mathcal{D}_D$ respectively, and u_1 and u_2 are two unlabeled examples, one from $\mathcal{D}_R$ and one from $\mathcal{D}_D$ (ordered randomly). **Left:** Attacker $\mathcal{M}_A$ targets only the trained Defender model $\mathcal{M}_D$. **Right:** $\mathcal{M}_A$ targets both $\mathcal{M}_D$ and its trainer $\mathcal{T}_D$.