

# Algorithm Selection for Optimal Multi-Agent Path Finding via Graph Embedding

**Abstract.** Multi-agent path finding (MAPF) is the problem of finding paths for multiple agents such that they do not collide. Finding optimal solutions to MAPF is NP-Hard, yet modern optimal solvers can scale to hundreds of agents and even thousands in some cases. Different solvers employ different approaches, and there is no single state of the art approach all problems. Furthermore, there are no clear, provable, guidelines for choosing when each optimal MAPF solver to use. Prior work employed Algorithm Selection (AS) techniques to learn such guidelines from past data. A major challenge when employing AS for choosing an optimal MAPF algorithm is how to encode the given MAPF problem. Prior work either used hand-crafted features, graph embedding of the shortest paths, or an image representation of problem. Each encoding is lossy, in the sense that it does not capture some aspect of the MAPF problem. We propose a graph encoding of the MAPF problem, and show how it can be used on-the-fly with a modern graph embedding algorithm called FEATHER. We also show how this encoding can be effectively joined with existing encodings, resulting in a novel AS method we call MAPF Algorithm selection via Graph embedding (*MAG*). An extensive experimental evaluation of *MAG* on several MAPF algorithm selection tasks reveals that it outperforms existing methods significantly.

## 1 Introduction

Multi-Agent Pathfinding (MAPF) is the problem of finding paths for a group of agents, moving each agent from its initial location to a designated target location. The main constraint in MAPF is that the agents must not collide. MAPF has received significant interest recently in the scientific community and in industry, as it has applications in robotics [28] and automated warehouses [29]. Finding an optimal solution to MAPF with respect to various optimization criteria, is known to be NP Hard [26, 31]. Nevertheless, a range of practical algorithms that guarantee optimality exists [5, 25]. It has been shown that these algorithms are able to find optimal solutions to MAPF problems with more than 100 agents with less than one minute of runtime [14].

Different optimal MAPF algorithms employ different problem solving techniques. For example, EPEA\* [8] employs a heuristic search technique, BCP [13] uses optimization techniques, and SAT-MDD [27] compiles MAPF to Boolean Satisfiability. Correspondingly, different solvers work best for different MAPF problems, and no algorithm has emerged to dominate all others. The variance in performance can be great, where some algorithms perform poorly on some instances but significantly outperform all others on other instances. On a recently performed extensive comparison of 5 optimal

MAPF algorithms [11], it was shown that even the least effective algorithm had some grids in which it was able to solve problems with 4 times more agents than all other evaluated algorithms. Thus, developing methods for selecting the best optimal MAPF search algorithm for a given problem is a worthwhile endeavor.

The problem of determining which algorithm from a given set of algorithms is expected to perform best on a given problem instance is known as the Algorithm Selection (AS) problem [18, 12].<sup>1</sup> Several recent works have developed AS techniques for optimal MAPF [10, 17, 1]. They used supervised learning to train a classifier that chooses the best optimal MAPF algorithm for a given MAPF problem, from a portfolio of optimal MAPF algorithms that include EPEA\* [8], ICTS [21], SAT-MDD [27], CBSH [6], and Lazy CBS [7]. A key challenge faced by all prior work on AS for optimal MAPF is how to encode a given MAPF problem, as an input to the supervised learner. Prior works proposed hand-crafted MAPF-related features, casting the MAPF problem as an image [1], and graph embedding of a subgraph containing the shortest paths from start to goal. Neither of these encodings completely captures the encoded MAPF problem.

To this end, we propose two contributions. The first contribution is a different encoding of the MAPF problem called *FG2V*, that fully utilizes the power of graph embedding algorithms by encoding the entire graph, augmented with artificial edges marking the start and goal vertex of every agent. This embedding is done using FEATHER [19], a modern graph embedding algorithm. The resulting embedding yields superior results in most cases, but not always. The second contribution is a simple method for integrating multiple encodings, which enables a more comprehensive encoding of the problem. This method can be used with different embeddings in a seamless manner. The resulting AS method is called *MAG*. Our third contribution is a comprehensive evaluation of *MAG* on a standard benchmark over three different AS tasks, which differ in how similar are the train and test problems. Our results show that *MAG* is superior to baseline methods, demonstrating the first practical AS method for MAPF that utilizes graph embeddings.

## 2 Background

For completeness, we provide here a brief background on MAPF, AS for MAPF, and graph embedding.

---

<sup>1</sup> Technically, this problem is called the *per-instance AS* problem.

## 2.1 MAPF

A MAPF problem is defined by a tuple  $\langle k, G, s, t \rangle$  where  $k$  is the number of agents,  $G = (V; E)$  is an undirected graph,  $s : [1 \dots, k] \rightarrow V$  maps an agent to its source vertex, and  $t : [1 \dots, k] \rightarrow V$  maps an agent to its target vertex. Initially each agent is in its source vertex. In every time step, each agent either waits in its current vertex or moves to one of the vertices adjacent to it. A single-agent plan for agent  $i$  is a sequence of move/wait actions that moves  $i$  from  $s(i)$  to  $t(i)$ . A solution to a MAPF problem is a set of single-agent plans, one for each agent, such that the agents do not collide with each other.

In the classical MAPF problem, the cost of a solution is either the sum of actions in all single-agent plans (also known as the sum of costs) or the number of actions in the longest single-agent plan (also known as makespan). MAPF extensions that consider non-unit action costs, large agents, and continuous time have been explored [3, 2, 15]. In this work we focus on classical MAPF.

Optimal MAPF algorithms are algorithms that are guaranteed to return optimal solutions according to a predefined solution cost function. Two commonly studied MAPF solution cost functions are the *sum of costs* (SOC) and *makespan*. SOC is the sum of actions (move or wait) performed by all agents until all agents reach their goals. Makespan is the maximal number of actions each agent makes until all agents reach their goal. While our approach is agnostic to the chosen cost function, we used considered SOC in our experiments and unless stated otherwise consider a solution to be optimal if it minimizes SOC.

Different optimal algorithms have been proposed for solving classical MAPF problems. Prime examples are EPEA\* [8], SAT-MDD [26], CBS [20, 6] and its many extensions, BCP [13], Lazy CBS [7], and ICTS [21]. These algorithms apply a range of techniques: some use heuristic search on dedicated state spaces, other compile the problem to Boolean Satisfiability (SAT) and call an off-the-shelf SAT solver, while others borrow ideas from constraint programming. No algorithm fully dominates the other, and different problems are solved best with different algorithms. This raised the need for an automated way to select which optimal MAPF solver to use for a given problem.

## 2.2 AS for MAPF

Previously proposed AS methods for MAPF followed a standard supervised learning paradigm, and differ mainly in the type of features they extract. Sigurdson et al. [22] and Ren et al. [17] mapped a given MAPF problem to an image and extracted image-based features with a Convolutional Neural Networks (CNN). Kaduri, Boyarski, and Stern [10] proposed a set of hand-crafted, MAPF-specific features, such as the number of agents divided by the number of unblocked cells in the grid. We refer to these features as the *KBS* features. Ren et al. also explored the potential of using features that are based on mapping the given MAPF problem to a graph and using a *graph embedding* method. While their exploration showed that graph embedding can provide useful features for MAPF AS, they do not provide a practical method to use it. In fact, they consider their method to be “not a deployable algorithm selector in any reasonable sense” [17], since it could not be used on any MAPF problem not observed during training. We overcome this limitation in our work.

Kaduri et al. [11] distinguished between three types of AS problems for MAPF: (1) in-grid AS, (2) in-grid-type, and (3) between-grid-type. In-grid AS means the train and test problems are all from

the same underlying graph. In-grid-type means the train and test problems are on different grids, but their grids are similar topologically. Between-grid-type means that graphs in the train and test problems are completely different, e.g., training on maze-like graphs and testing on graphs that represent road map in a city. Most prior work has focused on the in-grid AS problem, which is, of course, significantly easier.

## 2.3 Graph Embedding Algorithms

Node embedding methods are algorithms for encoding a node in a graph into a low-dimensional continuous vector [9]. Similarly, graph embedding methods are algorithms for encoding an entire graph  $G$  into a low-dimensional continuous vector [9], referred to as the *embedding* of  $G$ . Ideally, graphs with similar structure will have embedding that are close in terms of their Euclidean distance. Graph embeddings have proven to be useful features for various machine learning tasks such as classification [30].

Graph2Vec [16] is a neural graph embedding algorithm that accepts a set of graphs and outputs an embedding for each graph by analyzing local neighborhood of their nodes. It runs stochastic gradient descent to optimize a loss function that ensures similar graphs in the input set of graphs will have a similar embedding. Notably, it cannot be effectively used on other graphs without re-training, and thus its usefulness for our purposes is limited. FEATHER [19] is a recently proposed algorithm that serves as both a node embedding and a graph embedding algorithm. Its embedding is based on the likelihood of reaching each node in a random walk. Unlike Graph2Vec, it does not require an optimization step, and can reasonably used to embed a single graph. The embeddings created by FEATHER have shown to be effective in node-level and graph-level machine learning tasks, such as classifying fake users in a social network and link prediction. FEATHER has an additional positive property that it describes isomorphic graphs with the same representation and exhibits robustness to data corruption.

## 3 Method

In this section, we describe our AS method for choosing optimal classical MAPF algorithms called MAPF Algorithm selection via Graph embedding (MAG). MAG works as follows. First, it encodes the given MAPF problem as a graph. Then, it creates an embedding of the resulting graph with FEATHER. The resulting vector is added to a set of previously proposed MAPF-specific features extracted from the given MAPF problem, and used to solve our AS problem using supervised machine learning. Next, we describe each of these steps in more detail.

### 3.1 Encoding MAPF as a Graph

We consider two ways to encode a MAPF problem  $\Pi = \langle k, G, s, t \rangle$  as a graph. The first encoding method, called G2V, was previously proposed by Ren et al. [17]. G2V encode  $\Pi$  as the subgraph of  $G$  that includes the shortest paths for each agent and the nodes adjacent to them. In more details, G2V computes for each agent the shortest path from its source to its target while ignoring all other agents. Then, it joins these paths to a single graph and adds links if any between adjacent path’s nodes in original MAPF problem. Note that the resulting graph may be disconnected.

The benefit of G2V is that the resulting graphs are significantly smaller than  $G$ . However, these graphs lose information: they do not

consider regions on  $G$  that are not on the agents shortest paths. These regions may be important to consider when the corresponding MAPF problem is particularly difficult and agents must move away from their shortest paths. To address this limitation of  $G2V$  we propose  $FullG2V$  ( $FG2V$ ), which uses the entire graph  $G$  to encode the given MAPF problem. To include details about the MAPF problem beyond  $G$ , the graph outputted by  $FG2V$  includes an artificial edge for every agent between its source and target. Formally, for a MAPF problem  $\langle k, G = (V, E), s, t \rangle$   $FG2V$  outputs the graph  $G' = (V', E')$  where  $V' = V$  and  $E' = E \cup \{(s(i), t(i))\}_{i=1}^k$ . Figure 1 illustrates  $G2V$  and  $FG2V$  on a simple MAPF problem.




Figure 1. An example of the  $G2V$  and  $FG2V$  encoding methods.

### 3.2 Embedding the Graph

To embed the encoded graph into a vector space, we used the FEATHER algorithm [19]. FEATHER creates a graph embedding by firsts embedding each of the graph nodes and then pooling the resulting vectors to a single vector of size 500. Unlike other graph embedding techniques, it does not require a-priori training. Consequently, the features extracted using FEATHER can be extracted and used in a meaningful way even for graphs created for MAPF problems that are not in the training set. This is key to allowing graph embedding features to be used for optimal MAPF AS methods.

Note that the default pooling method in FEATHER “mean”. This means the graph embedding is created by taking the mean over its constituent node embeddings. We observed that using mean pooling did not perform well in our context, i.e., led to poor classification results when used as features. The reason for this is that mean pooling diminishes the impact of the artificial edges added between the source and target of each agents. Thus, MAPF problems on the same graph yielded very similar embeddings. To overcome this, we configured FEATHER to use “max” pooling, which emphasizes small differences between graphs created from MAPF problems on the same grid. This yielded significantly better results when training the AS classification model.

### 3.3 Feature Extraction and Learning

For a given MAPF problem  $\Pi$ ,  $MAG$  uses  $G2V$  and  $FG2V$  to create two graphs  $G_{G2V}$  and  $G_{FG2V}$  that encoding  $\Pi$ . Then, it creates two graph embeddings  $v_{G2V}$  and  $v_{FG2V}$  by applying FEATHER on these graphs. It also creates a vector  $v_{KBS}$  by extracting all the MAPF-specific features proposed by Kaduri et al. [10]. Finally,  $MAG$  concatenates  $v_{G2V}$ ,  $v_{FG2V}$ , and  $v_{KBS}$ . The resulting vector is the features  $MAG$  uses for training. The training process itself is a straightforward multi-class supervised learning process, where every instance is a MAPF problem and the label is the fastest algorithm for that problem within our portfolio of algorithms. Figure 2 illustrates

the feature extraction and training process. Note that more sophisticated approaches to AS exists in other domains [12], e.g., methods that involve runtime predictions. This is left for future work.




Figure 2. Diagram of the  $MAG$  feature extraction and training process.

## 4 Experimental Results

In this section, we present an experimental evaluation of  $MAG$  on a standard publicly available grid-based MAPF benchmark [23]. This benchmark contains 33 grids arranged into seven *grid types*: video games (denoted as “game” grids), city maps (“city”), maze-like grids (“maze”), grids arranged as rooms with narrow doors between them (“room”), open grids (“empty”), open grids with randomly placed obstacles (“random”), and grids that are inspired by the structure of warehouses (warehouse). Figure 3 shows an example grid from each type.<sup>2</sup> The benchmark includes *scenario files* for each grid. Each scenario file contains source and target locations for up to 1,000 agents where possible. The scenario files of each grid are grouped into two sets. In the first set of scenario files, denoted *Random*, the agents source and target locations are located purely randomly. In the second set of scenario files, denoted *Even*, the agents’ source and target locations are evenly distributed in buckets of 10 agents according to their distance. Only the scenarios from the Even set were used, as they represent a more diverse set of MAPF problems.

### 4.1 Experimental Setup

We performed three sets of experiments, for each of the AS problem setups defined by Kaduri et al. [11]: in-grid, in-grid-type, and between-grid-type. To train and evaluate  $MAG$  in each setup, we used the publicly available dataset of Kaduri et al. [11], which includes results for running a set of optimal MAPF solvers of the entire MAPF benchmark mentioned above. Specifically, results for the following optimal MAPF solvers are available in this dataset: ICTS [21], EPEA\* [8], SAT-MDD [27], CBSH [6], and Lazy CBS [7].


#### 4.1.1 Baselines

We compared  $MAG$  against two baselines:

- **KBS** The AS method by Kaduri, Boyarski and Stern [10], which uses only their hand crafted features.
- **G2V** The AS method described by Ren et al.[17], which uses only the  $G2V$  encoding. To extract features from the  $G2V$  graph, we used the FEATHER graph embedding.<sup>3</sup>

<sup>2</sup> The images were taken from the MovingAI repository [24], which hosts the grid MAPF benchmark we used [23].

<sup>3</sup> This differs from Ren et al., who used Graph2Vec. As explained earlier, Graph2Vec is not a practical method for our problem, since it requires knowing a-priori the graphs to embed.



**Figure 3.** An example grid from each of the grid types in our benchmark, From left to right: empty, random, warehouse, game, city, maze, and room.

In addition, we performed an ablation study for *MAG*, and report on results for AS methods that use different subsets of features. Namely, *FG2V* + *G2V* *KBS* + *FG2V* *KBS* + *G2V* and *FG2V*. Note that *KBS* + *G2V* + *FG2V* is exactly *MAG*. For training, we used XGBoost [4], a well-known supervised learning algorithm. Preliminary experiments with other learning algorithms, such as Logistic Regression and Random Forest, yielded weaker results. The hyper parameters of XGBoost were tuned by performing a 4-fold cross validation over the training set, for each AS setup.

#### 4.1.2 Metrics

The main metrics used in prior work on AS for MAPF are:

- **Accuracy (Acc).** Ratio of problems where the AS method returned the fastest MAPF algorithm.
- **Coverage (Cov).** Ratio of MAPF problems solved under a time limit of 5 minutes.<sup>4</sup>
- **Runtime (RT).** Average run-time in minutes to solve a single MAPF problems with the selected MAPF solver.<sup>5</sup>

Note that since all solvers are optimal MAPF solvers, all solvers return solutions of exactly the same cost. Thus, comparing solution quality is redundant.

To provide context for our results, we also report on the results of an *Oracle*, which always selects the fastest algorithm for every MAPF problem. No practical AS method can perform better than Oracle, which has an accuracy and coverage of 1.0, and the smallest possible runtime. Based on the runtime of Oracle, we also report for every AS method on the average runtime it required over A final metric, we report for every AS method the average percentage of Oracle runtime required for the selected MAPF solver beyond the runtime required by the algorithm selected by Oracle. We call this metric the *average regret*, or briefly %Rg. The average regret of Oracle is by definition zero, and better AS methods will have lower average regret values.

## 4.2 In-Grid Results

Table 4.2 presents the results for the in-grid AS setup experiments. The rows correspond to different AS methods, and the columns are the metrics defined earlier. The column groups “All” and “Avg” provide a slightly different way to aggregate the results over all test problems. The results under the “All” columns are averages over all test problems, regardless of their grids and grid types. This is how most prior work on AS for MAPF aggregated their results. The limitation of this aggregation is that some grids in the benchmark are smaller than others, and has fewer MAPF problems defined for them. The results under the “Avg” columns are averages of averages, where the

<sup>4</sup> This time limit is common in the Optimal MAPF literature.

<sup>5</sup> We considered cases where the selected MAPF solver could not solve the problem within our 5 minutes time limit as having a runtime of 5 minutes. The same has been done in prior work on MAPF AS [10, 17].

Metric	All			Avg			
	Acc	Cov	RT	Acc	Cov	RT	%Rg
KBS	0.83	0.98	0.549	0.88	<b>0.99</b>	0.41	12.4
G2V	0.81	0.97	0.589	0.87	0.98	0.45	25.4
<i>MAG</i>	<b>0.85</b>	<b>0.99</b>	<b>0.514</b>	<b>0.89</b>	<b>0.99</b>	<b>0.40</b>	<b>9.0</b>
Oracle	1.00	1.00	0.436	1.00	1.00	0.36	0.0
FG2V+G2V	0.84	0.98	0.531	0.89	0.99	0.41	11.5
KBS+G2V	0.84	0.98	0.532	0.89	0.99	0.40	9.5
KBS+FG2V	0.85	0.99	0.513	0.89	0.99	0.39	8.1
FG2V	0.84	0.98	0.534	0.88	0.99	0.41	11.8

**Table 1.** Results for the In-Grid AS setup, averaged over all test problems.

results of each grid type are averaged separately and only the resulting averages are averaged. This mitigates unwanted to bias stemming from the number of problems in each grid type.

Consider first the results for *MAG* and our two main baselines, *KBS* and *G2V*. For each metric, we highlighted the best results among these AS methods in bold. As the results clearly show, *MAG* is either on par or better than these baselines on all metrics. For example, the accuracy of *MAG* is 0.85 in “All” while it is 0.83 and 0.81 for *KBS* and *G2V* respectively. The advantage of *MAG* over *G2V* is more significant, and more modest compared to *KBS*. Still advantage is significant, especially in terms of the average regret, which is approximately 25% smaller than *KBS* (9.0 vs. 12.4).

When analyzing the ablation study baselines (last 4 rows in Table 4.2, we see very similar results to *MAG* where there is a slight advantage for using the *KBS* hand-crafted MAPF-specific features together with graph embedding features (either *KBS* + *G2V* or *KBS* + *FG2V*). This is expected, as more diverse set of features is expected to be more beneficial.

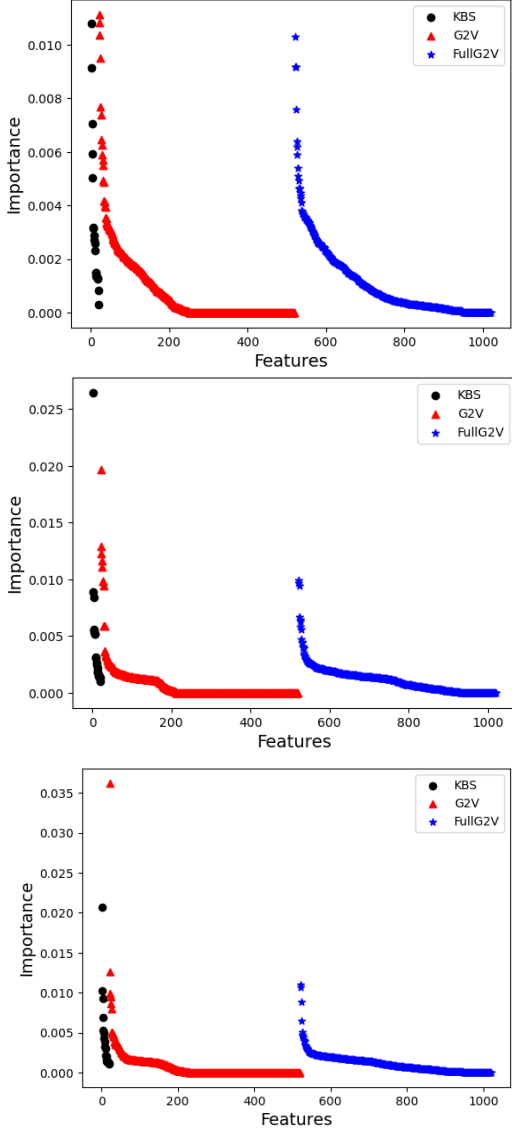
## 4.3 In-Grid-Type and Between-Grid-Type Results

Metric	In-Grid-Type				Between-Grid			
	Acc	Cov	RT	% Rg	Acc	Cov	RT	% Rg
KGS	0.69	0.93	0.86	78.6	<b>0.63</b>	0.87	0.99	223.1
G2V	0.67	0.92	0.91	91.8	0.61	0.86	1.04	220.7
<i>MAG</i>	<b>0.71</b>	<b>0.94</b>	<b>0.80</b>	<b>67.9</b>	<b>0.63</b>	<b>0.89</b>	<b>0.93</b>	<b>180.0</b>
Oracle	1.00	1.00	0.48	0.00	1.00	1.00	0.36	0.00
FG2V+G2V	0.69	0.93	0.84	75.8	0.62	0.88	0.97	195.8
KGS+G2V	0.70	0.94	0.81	70.5	0.64	0.89	0.93	192.5
KGS+G2V	0.70	0.94	0.82	71.6	0.65	0.88	0.94	181.8
FG2V	0.66	0.91	0.90	87.9	0.62	0.88	0.95	176.0

**Table 2.** In-grid-type (left) and Between-grid-type (right) results.

The first 5 columns (left-to-right) in Table 4.3 presents the results for the in-grid-type AS setup experiments. Similar to Table 4.2, the rows are different AS methods and the columns are different metrics, corresponding to the “Avg” column family. We highlighted in bold the AS method, among *MAG* and our two baselines, that yielded the





**Figure 4.** Feature importance for the XGBoost model created for *MAG* in the in-grid (top), in-grid-type (middle), and between-grid (bottom) setups.

best results in each metric. As in the in-grid experiments, the advantage of *MAG* over the baselines is clear in all metrics. For example, its average regret is 67.9 while it is 78.60 and 91.80 for *KBS* and *G2V* respectively.

The ablation study results show that here too, combining the hand-crafted features of *KBS* with either type of graph embedding provides the biggest performance improvement. For example, *KBS* with either *G2V* or *FG2V* yields 0.70 accuracy and runtime of 0.82 or less while *FG2V* alone or even with *G2V* yielded lower accuracy and a higher runtime. It is worth comparing the average regret results here and in the in-grid experiments. While the regret of *MAG* here is 67.90 it is only 9.0 in the in-grid results. This highlights that in-grid-type AS is a significantly harder task than the in-grid AS, since it requires generalizing from different grids (although from the same type).

The rightmost 4 columns in Table 4.3 present the results for the between-grid-type AS setup experiments. The first trend we observe is that the overall results for all algorithms is significantly worse compared to all other AS setups (in-grid and in-grid-type). For example,

the average regret of *MAG* in the in-grid-type results is 67.9% but it is 180.0% in between-grid results. Similarly, *MAG* accuracy dropped from 0.71 to 0.63. Recall that the regret and accuracy of *MAG* in the in-grid was 9.0% and 0.89, respectively. These differences are expected, since in the between-grid experiments, the training set did not include any grid of the tested type, which makes the classification problem significantly harder.

In terms of the comparison with our baselines, the general trend we observed so far continues in the between-grid setup: *MAG* is either on par or better than the baselines in all metrics. For example, its average runtime and regret are 0.93 and 180 while it is 0.99 and 220.7 for the next best baseline, respectively. The ablation study results are less conclusive in this setup. In terms of accuracy, we still see the benefit in combining *KBS* features with graph embedding features over only using graph embedding features. However, the lowest average regret is achieved when only using *MAG*. In fact, some subsets of *MAG* features actually outperform *MAG* on some metrics. For example, using only *FG2V* yields lower accuracy than the full *MAG* but a slightly lower average regret (176 vs. 180). However, these differences are relatively small.

#### 4.4 Grid Types Analysis

Table 4.3 provides a deeper insight into the results of our in-grid, in-grid-type, and between-grid experiments. Here, the results — accuracy, coverage, and regret — are grouped by grid types. For example, the results in the row “City” show the average results over test problems that are on grids of type “City”. The rows correspond to the test grid type, and the columns correspond to the AS setup and evaluated algorithm.

While *MAG* is still, in general, the best-performing algorithm, in some grid types and metrics it is not, especially in the between grid setup. This is most evident in the results for Maze grids, where the accuracy, coverage, and regret of *KBS* — 0.56, 0.71, and 398.8%, respectively — are better than the corresponding results of *MAG* which were 0.47, 0.66, and 483.6%. A possible explanation for these results is that the Maze grid are significantly different from the other grid types in their topology, where wrong path finding choices can have a large impact on performance and are harder to “recover” from, as they require backtracking to other choices in the maze.

Interestingly, *MAG* performs well on City grids in the in-grid AS setup. This suggests it is able to learn how to act in such grids, if they are given to it for training. The in-grid results for Maze grids provide another interesting insight when comparing the results of *MAG* and *G2V*. As can be seen, the difference in this grid type between *MAG* and *G2V* is most pronounced: the accuracy, coverage, and regret of *G2V* are 0.79, 0.93, and 119.3, respectively, while they are 0.86, 0.99, and 20.6 for *MAG*. The poor performance of *G2V* on Maze grids is understandable: in such grids agents often follow paths that are significantly different from their shortest paths to avoid collisions. Such paths are not encoded by *G2V*. In contrast, the *FG2V* works quite well on these grids, almost as well as *MAG*.

#### 4.5 Feature Importance

Figure 4 plots a feature importance analysis performed on the prediction models created by *MAG* for each of the AS setup. The features are listed on the x-axis, where the first 20 features are the *KBS* features, the next 500 features are the *G2V* features, and the last 500 features are the *FG2V* features. Within each feature family (*KBS*, *G2V*, and *FG2V*), the features are sorted by their importance.

Grid-type	AS Setup Metric	In-Grid			In-Grid-Type			Between-Grid		
		KBS	G2V	MAG	KBS	G2V	MAG	KBS	G2V	MAG
Empty	Acc	0.88	0.89	<b>0.90</b>	0.67	0.68	<b>0.71</b>	0.81	0.78	<b>0.82</b>
	Cov	<b>1.00</b>	<b>1.00</b>	<b>1.00</b>	0.97	<b>1.00</b>	0.99	<b>1.00</b>	0.99	<b>1.00</b>
	%Rg	0.9	<b>0.1</b>	0.8	3.5	<b>1.0</b>	1.8	9.5	17.7	<b>2.5</b>
Random	Acc	<b>0.91</b>	0.90	<b>0.91</b>	<b>0.84</b>	0.79	0.82	<b>0.78</b>	0.71	0.73
	Cov	<b>1.00</b>	<b>1.00</b>	<b>1.00</b>	<b>1.00</b>	0.99	<b>1.00</b>	<b>1.00</b>	0.99	<b>1.00</b>
	%Rg	3.1	4.1	<b>2.4</b>	1.4	2.3	<b>1.1</b>	7.2	54.8	<b>4.0</b>
Warehouse	Acc	0.84	0.84	<b>0.86</b>	0.64	0.64	<b>0.67</b>	0.67	0.62	<b>0.68</b>
	Cov	0.99	0.99	<b>1.00</b>	0.94	0.94	<b>0.96</b>	0.96	0.88	<b>0.98</b>
	%Rg	7.9	9.4	<b>4.7</b>	1.4	1.6	<b>1.3</b>	32.9	78.1	<b>27.3</b>
Game	Acc	0.91	0.91	<b>0.92</b>	0.77	0.72	<b>0.78</b>	<b>0.74</b>	0.72	0.65
	Cov	0.98	0.98	<b>0.99</b>	<b>0.92</b>	0.86	<b>0.92</b>	0.88	<b>0.89</b>	<b>0.89</b>
	%Rg	19.2	21.0	<b>18.0</b>	<b>1.4</b>	2.3	1.7	122.0	113.1	<b>122.0</b>
City	Acc	0.90	0.89	<b>0.92</b>	0.58	0.57	<b>0.61</b>	0.53	0.53	<b>0.58</b>
	Cov	<b>0.99</b>	<b>0.99</b>	<b>0.99</b>	0.81	<b>0.83</b>	<b>0.83</b>	<b>0.90</b>	0.86	0.88
	%Rg	10.1	12.4	<b>7.1</b>	3.1	3.0	<b>2.9</b>	<b>136.2</b>	155.6	146.3
Maze	Acc	0.85	0.79	<b>0.86</b>	0.40	0.44	<b>0.50</b>	<b>0.56</b>	0.45	0.47
	Cov	0.98	0.93	<b>0.99</b>	<b>0.66</b>	0.62	<b>0.66</b>	<b>0.71</b>	0.64	0.66
	%Rg	31.9	119.3	<b>20.6</b>	<b>8.2</b>	9.6	8.4	<b>398.8</b>	511.4	483.6
Room	Acc	<b>0.93</b>	0.90	0.92	0.63	0.50	<b>0.66</b>	0.32	0.43	<b>0.49</b>
	Cov	<b>1.00</b>	<b>1.00</b>	<b>1.00</b>	0.91	0.83	<b>0.94</b>	0.65	0.75	<b>0.81</b>
	%Rg	0.2	2.1	<b>0.1</b>	2.9	5.0	<b>2.2</b>	855.6	614.5	<b>475.0</b>

**Table 3.** Results for all AS setups, grouped by test grid type.

Our main observation is that each feature family has some features that are important, and some features whose importance. This suggests there is useful information in all types of features, which corresponds to the successful results of *MAG*. Another observation is that each feature family includes features whose importance is close to zero. This suggests that applying feature selection method based on weights may be effective. When comparing the different AS setups, it seems that as the AS setup becomes more challenging, the set of features that are important is reduced, and the importance of more features is close to zero. This may suggest that additional types of features may be needed to get better results for the harder AS setups.

## 5 Conclusion and Future Work

We proposed *MAG*, the first practical approach to optimal MAPF algorithm selection based on graph embedding. *MAG* uses two encodings of the MAPF problem as a graph, one that encodes the entire graph and one that encodes only the shortest paths and their immediate vicinity. To work efficiently on new MAPF problems, *MAG* utilizes a modern graph embedding algorithm that does not need a-priori training. *MAG* also uses hand-crafted MAPF-specific features, as suggested by prior work. The combination of graph-embedding features and hand-crafted features leads to strong state of the art AS for optimal MAPF. In an extensive set of experiments on a standard benchmark, we showed that *MAG* significantly outperforms existing baselines almost always. Our results also highlight that the between-grid AS setup is particularly challenging, and can be the focus of future work. Another important future work is to combine our graph-embedding based AS model with image-based AS models, such as MAPFASTER [1].

## References

- [1] Jean-Marc Alkazzi, Anthony Rizk, Michel Salomon, and Abdallah Makhoul, ‘Mapfaster: A faster and simpler take on multi-agent path finding algorithm selection’, in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pp. 10088–10093, (2022).
- [2] Anton Andreychuk, Konstantin Yakovlev, Pavel Surynek, Dor Atzmon, and Roni Stern, ‘Multi-agent pathfinding with continuous time’, *Artificial Intelligence*, (2022).
- [3] Dor Atzmon, Yonathan Zax, Einat Kivity, Lidor Avitan, Jonathan Morag, and Ariel Felner, ‘Generalizing multi-agent path finding for heterogeneous agents’, in *International Symposium on Combinatorial Search (SoCS)*, volume 11, pp. 101–105, (2020).
- [4] Tianqi Chen and Carlos Guestrin, ‘Xgboost: A scalable tree boosting system’, in *Proceedings of the 22nd acm sigkdd international conference on knowledge discovery and data mining*, pp. 785–794, (2016).
- [5] Paul Spirakis, Daniel Kornhauser, Gary Miller, ‘Coordinating pebble motion on graphs, the diameter of permutation groups, and applications’, in *FOCS*, pp. 241–250, (1984).
- [6] Ariel Felner, Jiaoyang Li, Eli Boyarski, Hang Ma, Liron Cohen, TK Satish Kumar, and Sven Koenig, ‘Adding heuristics to conflict-based search for multi-agent path finding’, in *International Conference on Automated Planning and Scheduling*, volume 28, pp. 83–87, (2018).
- [7] Graeme Gange, Daniel Harabor, and Peter J Stuckey, ‘Lazy cbs: implicit conflict-based search using lazy clause generation’, in *International Conference on Automated Planning and Scheduling*, volume 29, pp. 155–162, (2019).
- [8] Meir Goldenberg, Ariel Felner, Roni Stern, Guni Sharon, Nathan R. Sturtevant, Robert C. Holte, and Jonathan Schaeffer, ‘Enhanced partial expansion A\*’, *J. Artif. Intell. Res.*, **50**, 141–187, (2014).
- [9] Palash Goyal and Emilio Ferrara, ‘Graph embedding techniques, applications, and performance: A survey’, *Knowledge-Based Systems*, **151**, 78–94, (2018).
- [10] Omri Kaduri, Eli Boyarski, and Roni Stern, ‘Algorithm selection for optimal multi-agent pathfinding’, in *Proceedings of the International Conference on Automated Planning and Scheduling*, volume 30, pp. 161–165, (2020).
- [11] Omri Kaduri, Eli Boyarski, and Roni Stern, ‘Experimental evaluation of classical multi agent path finding algorithms’, in *Proceedings of the In-*

- ternational Symposium on Combinatorial Search, volume 12, pp. 126–130, (2021).
- [12] Pascal Kerschke, Holger H Hoos, Frank Neumann, and Heike Trautmann, ‘Automated algorithm selection: Survey and perspectives’, *Evolutionary computation*, **27**(1), 3–45, (2019).
  - [13] Edward Lam, Pierre Le Bodic, Daniel Harabor, and Peter J Stuckey, ‘Branch-and-cut-and-price for multi-agent path finding’, *Computers & Operations Research*, **144**, 105809, (2022).
  - [14] Jiaoyang Li, Daniel Harabor, Peter J Stuckey, Hang Ma, Graeme Gange, and Sven Koenig, ‘Pairwise symmetry reasoning for multi-agent path finding search’, *Artificial Intelligence*, **301**, 103574, (2021).
  - [15] Jiaoyang Li, Pavel Surynek, Ariel Felner, Hang Ma, TK Satish Kumar, and Sven Koenig, ‘Multi-agent path finding for large agents’, in *AAAI Conference on Artificial Intelligence*, volume 33, pp. 7627–7634, (2019).
  - [16] Annamalai Narayanan, Mahinthan Chandramohan, Rajasekar Venkatesan, Lihui Chen, Yang Liu, and Shantanu Jaiswal, ‘Graph2Vec: Learning distributed representations of graphs’, *arXiv:1707.05005*, (2017).
  - [17] Jingyao Ren, Vikraman Sathiyarayanan, Eric Ewing, Baskin Senbaslar, and Nora Ayanian, ‘MAPFAST: A deep algorithm selector for multi agent path finding using shortest path embeddings’, in *International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, eds., Frank Dignum, Alessio Lomuscio, Ulle Endriss, and Ann Nowé, pp. 1055–1063, (2021).
  - [18] John R Rice, ‘The algorithm selection problem’, in *Advances in computers*, volume 15, 65–118, (1976).
  - [19] Benedek Rozemberczki and Rik Sarkar, ‘Characteristic functions on graphs: Birds of a feather, from statistical descriptors to parametric models’, in *ACM international conference on information & knowledge management*, pp. 1325–1334, (2020).
  - [20] Guni Sharon, Roni Stern, Ariel Felner, and Nathan R Sturtevant, ‘Conflict-based search for optimal multi-agent pathfinding’, *Artificial Intelligence*, **219**, 40–66, (2015).
  - [21] Guni Sharon, Roni Stern, Meir Goldenberg, and Ariel Felner, ‘The increasing cost tree search for optimal multi-agent pathfinding’, *Artificial intelligence*, **195**, 470–495, (2013).
  - [22] Devon Sigurdson, Vadim Bulitko, Sven Koenig, Carlos Hernandez, and William Yeoh, ‘Automatic algorithm selection in multi-agent pathfinding’, *arXiv:1906.03992*, (2019).
  - [23] Roni Stern, Nathan Sturtevant, Ariel Felner, Sven Koenig, Hang Ma, Thayne Walker, Jiaoyang Li, Dor Atzmon, Liron Cohen, TK Kumar, et al., ‘Multi-agent pathfinding: Definitions, variants, and benchmarks’, in *International Symposium on Combinatorial Search*, volume 10, pp. 151–158, (2019).
  - [24] N. Sturtevant, ‘Benchmarks for grid-based pathfinding’, *Transactions on Computational Intelligence and AI in Games*, **4**(2), 144 – 148, (2012).
  - [25] Pavel Surynek, ‘A novel approach to path planning for multiple robots in bi-connected graphs’, in *IEEE International Conference on Robotics and Automation*, pp. 3613–3619, (2009).
  - [26] Pavel Surynek, ‘An optimization variant of multi-robot path planning is intractable’, in *AAAI*, (2010).
  - [27] Pavel Surynek, Ariel Felner, Roni Stern, and Eli Boyarski, ‘Efficient sat approach to multi-agent path finding under the sum of costs objective’, in *European conference on artificial intelligence*, pp. 810–818, (2016).
  - [28] Manuela M Veloso, Joydeep Biswas, Brian Coltin, and Stephanie Rosenthal, ‘Cobots: Robust symbiotic autonomous mobile service robots’, in *International Joint Conference on Artificial Intelligence (IJCAI)*, (2015).
  - [29] Peter R Wurman, Raffaello D’Andrea, and Mick Mountz, ‘Coordinating hundreds of cooperative, autonomous vehicles in warehouses’, *AI magazine*, **29**(1), 9–9, (2008).
  - [30] Renchun You, Zhiyao Guo, Lei Cui, Xiang Long, Yingze Bao, and Shilei Wen, ‘Cross-modality attention with semantic graph embedding for multi-label classification’, in *AAAI conference on artificial intelligence*, volume 34, pp. 12709–12716, (2020).
  - [31] Jingjin Yu and Steven M. LaValle, ‘Structure and intractability of optimal multi-robot path planning on graphs’, in *AAAI*, (2013).