

Bounded Suboptimal Game Tree Search

Dor Atzmon and Roni Stern
Ben Gurion University of the Negev
Be'er Sheva, Israel

Abdallah Saffidine
Australian National University
Canberra, Australia

Abstract

.....

Bounded Alpha-Beta with multiplicative bound (BAB/m)

A bounded-suboptimal game tree search algorithm with multiplicative bound is defined as a game tree search algorithm that accepts $\gamma \geq 0$ as input and outputs a solution with suboptimality at most $\gamma \cdot V$ from the true minimax value V , where γ is a user-provided parameter. In our case, to adjust BAB to this type of bound, both L and U must be inside these bounds as they are part of the solution, thus: $V - \gamma \cdot V \leq L \leq U \leq V + \gamma \cdot V$, or $1 - \gamma \leq \frac{L}{V} \leq \frac{U}{V} \leq 1 + \gamma$. Algorithm 1 presents the changes needed to be done in BAB in order to return a bounded solution.

....

Worst / best case of BAB

Game plying with BAB

Improvements and comparison (null window search, history heuristic, and more)

Choosing a Bounded-Suboptimal Action

BAB (Algorithm 1) returns a range $[L(n_1), U(n_1)]$ for a given game tree and ϵ . As proven in Theorem ??, the minimax value is in that range, and the size of the range is at most ϵ . Thus, any value chosen in that range can serve as a bounded-suboptimal solution for the given ϵ .

To obtain a concrete policy from the output of BAB the MAX player needs to choose the action that leads to the highest L values. This policy is guaranteed (on expectation) to provide at least a bounded-suboptimal minimax value, for any action the MIN player may choose. By contrast, choosing the action that leads to the highest U values may end up with a strategy whose expected value is smaller than the minimax value by more than ϵ .

Copyright © 2018, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

ϵ -equilibrium

The optimal policy for MAX and MIN players that use classical Alpha-Beta pruning is known to be a Nash equilibrium, that is, no player can gain by unilaterally deviating from its action. Obviously, a the policy proposed in "" section obtained by BAB in not necessarily a Nash equilibrium. An ϵ -equilibrium (Nisan et al. 2007) is a joint policy in which no player can gain more than ϵ by unilaterally deviating from his strategy. Proving this conjecture is a topic for future research.

Theorem 1. For a game tree G and an ϵ bounded solution, playing the policy proposed in "" results in ϵ -equilibrium

Proof. Assume by the policy in "" child c' is selected.

Case #1: n is a MAX node. By theorem ??, for each child of n

$$U(c) - L(c) \leq \epsilon \quad (1)$$

And

$$V(c) \leq U(c) \quad (2)$$

By the policy in ""

$$\max_{c \in C(n)} L(c) = L(c') \leq V(c') \quad (3)$$

Using (2) and (3) in (1) we get that for each child c of n

$$V(c) - V(c') \leq \epsilon \quad (4)$$

And thus playing the policy proposed in "" results in ϵ -equilibrium \square

Related Work

....

Acknowledgements

This research was supported by the Israel Ministry of Science, the Czech Ministry of Education, and by ISF grant #0/17 to Roni Stern.

Algorithm 1: WeightedAB

Input: n - a game tree node
Input: ϵ - an error margin

- 1 if n is a terminal node then
- 2 \lfloor return ($V(n)$, $V(n)$)
- 3 $L(n) \leftarrow v_{\min}$; $U(n) \leftarrow v_{\max}$
- 4 if n is a MAX node then $best_U \leftarrow v_{\min}$
- 5 else if n is a MIN node then $best_L \leftarrow v_{\max}$
- 6 if n is n_I then $\alpha(n) = v_{\min}$; $\beta(n) = v_{\max}$
- 7 else if $p(n)$ is a MAX node or a MIN node then
 $\alpha(n) = \alpha(p(n))$; $\beta(n) = \beta(p(n))$
- 8 else
- 9 $\alpha(n) \leftarrow \max(v_{\min}, \frac{\alpha(p(n)) - U(p(n))}{\pi(n)} + U(n))$
- 10 $\beta(n) \leftarrow \min(v_{\max}, \frac{\beta(p(n)) - L(p(n))}{\pi(n)} + L(n))$
- 11 foreach child c of n do
- 12 $(L(c), U(c)) \leftarrow \text{WeightedAB}(c, \epsilon)$
- 13 if n is a MAX node then
- 14 $best_U \leftarrow \max(best_U, U(c))$
- 15 $L(n) \leftarrow \max(L(n), L(c))$
- 16 else if n is a MIN node then
- 17 $best_L \leftarrow \min(best_L, L(c))$
- 18 $U(n) \leftarrow \min(U(n), U(c))$
- 19 else
- 20 $L(n) \leftarrow L(n) + \pi(c)(L(c) - v_{\min})$
- 21 $U(n) \leftarrow U(n) - \pi(c)(v_{\max} - U(c))$
- 22 $\alpha(n) \leftarrow \max(L(n), \alpha(n))$
- 23 $\beta(n) \leftarrow \min(U(n), \beta(n))$
- 24 $\epsilon \leftarrow w \cdot \min(|\alpha(n)|, |\beta(n)|)$
- 25 if $\beta(n) \leq \alpha(n) + \epsilon$ and c is not the last child
 of n then return ($L(n)$, $U(n)$)
- 26 if n is a MAX node then $U(n) \leftarrow best_U$
- 27 else if n is a MIN node then $L(n) \leftarrow best_L$
- 28 return ($L(n)$, $U(n)$)

References

Nisan, N.; Roughgarden, T.; Tardos, E.; and Vazirani, V. V. 2007. Algorithmic game theory. Cambridge university press.