

# A Large-scale Dataset for Argument Quality Ranking: Construction and Analysis

Shai Gretz\*, Roni Friedman\*, Edo Cohen-Karlik\*, Assaf Toledo\*, Dan Lahav,  
Ranit Aharonov and Noam Slonim

IBM Research

{avishaig,roni.friedman-melamed,noams}@il.ibm.com  
{edo.cohen,assaf.toledo,dan.lahav,ranit.aharonov}@ibm.com

## Abstract

Identifying the quality of free-text arguments has become an important task in the rapidly expanding field of computational argumentation. In this work, we explore the challenging task of argument quality ranking. To this end, we created a corpus of 30,497 arguments carefully annotated for point-wise quality, released as part of this work. To the best of our knowledge, this is the largest dataset annotated for point-wise argument quality, larger by a factor of five than previously released datasets. Moreover, we address the core issue of inducing a labeled score from crowd annotations by performing a comprehensive evaluation of different approaches to this problem. In addition, we analyze the quality dimensions that characterize this dataset. Finally, we present a neural method for argument quality ranking, which outperforms several baselines on our own dataset, as well as previous methods published for another dataset.

## 1 Introduction

Computational Argumentation is a rapidly emerging discipline within the Natural Language Processing community (?), dealing with various sub-tasks such as argument detection (?; ?), stance detection (?) and argument clustering (?).

Recently, IBM introduced *Project Debater*, the first AI system able to debate humans on complex topics. The system participated in a live debate against a world champion debater, and was able to mine arguments, use them for composing a speech supporting its side of the debate, and also rebut its human competitor.<sup>1</sup> The underlying technology is intended to enhance decision-making.

More recently, IBM also introduced *Speech by Crowd*, a service which supports the collection of free-text arguments from large audiences on debatable topics to generate meaningful narratives (?). An important sub-task of this service is automatic assessment of argument quality, which is the focus of the present work. Detecting argument quality is a prominent task due to its importance in automated decision making (?), argument search (?), and writing support (?).

Earlier work on assessing argument quality relied on a comparative *pair-wise* approach, aiming to identify the higher-quality argument within each pair of arguments (?; ?). Recently ? (?) proposed a *point-wise* argument quality prediction scheme, which scales linearly with data size. Correspondingly, here we focus on this paradigm which is clearly less demanding, especially in scenarios where many arguments should be considered.

A major contribution of this work is introducing a novel dataset of arguments, carefully annotated for point-wise quality, *IBM-ArgQ-Rank-30kArgs*, referred henceforth as *IBM-Rank-30k*. The dataset includes around 30k arguments, 5 times larger than the largest annotated point-wise data released to date (?). Similarly to ? (?), arguments were collected actively – as opposed to being extracted from debate portals (?) – with strict length limitations, accompanied by extensive quality control measures.

Although a continuous argument quality score seems natural, asking annotators to provide a continuous score per argument will probably introduce a subjective scale that varies from one annotator to another, hindering downstream analysis. Instead, we take a simplified approach, asking each annotator to answer a binary question per argument, indicating if its quality is satisfactory in a particular context. The question remains, of how to extract a continuous quality score out of the binary annotations provided by many annotators. ? (?) took a straightforward simple-average approach; here, we provide an extensive comparison between potential scoring functions, analyzing the differences between these models, and their impact on training and evaluating a learning algorithm.

While an exact definition of argument quality is potentially elusive, it seems clear that it is a function of various linguistic phenomena. As exemplified in Table 1, low quality can be manifested by dimensions such as bad grammar and low clarity (Row 3), or lack of impact and relevance (Row 4). In contrast, high quality arguments are typically clear, relevant, and with high impact (Rows 1 and 2). A model that aims to automatically infer argument quality should take such subtleties into account. A recent development in this context is that of deep contextual language models, such as ELMo (?) and BERT (?). Due to its bidirectional nature,

\* These authors equally contributed to this work.

Copyright © 2020, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

<sup>1</sup>For more details: <https://www.research.ibm.com/artificial-intelligence/project-debater/live/>

BERT achieves remarkable results when fine-tuned to different tasks without the need for specific modifications per task. As part of this work, we introduce various neural methods that exploit the value of BERT for our task. In particular, we suggest a model that outperforms several baselines on our data. Our experimental results further indicate that this method is either comparable to or outperforms state-of-the-art methods on previously released data.

Argument	Topic	Label
the interest rates are too high and trap people in debt	Payday loans should be banned	1
racial profiling unfairly targets minorities and the poor	We should end racial profiling	1
we should subsidize student loans for reach excelent education	We should subsidize student loans	0.05
i think the same as you, they should ban	Payday loans should be banned	0.09

Table 1: Examples of high (rows 1-2) and low (rows 3-4) quality arguments from the *IBM-Rank-30k* dataset. The label is a weighted aggregation of annotations.

The main contributions of our work are: (1) Introducing a carefully annotated argument quality dataset which is the largest of its kind; (2) Conducting extensive analysis of different approaches to induce a quality label from given binary annotations; (3) Proposing a BERT-based method to predict argument quality, and report the results of extensive experiments that convey the potential of this method.

## 2 Related Work

Assessing argument quality is a long standing challenge. For centuries there has been a multi-disciplinary effort to define and research aspects of quality in argumentation (??; ?). A core issue in the field is the presumed subjectivity of the task at hand. There have been several practical and theoretical approaches on how to overcome the supposed lack of objectivity.

? (?) approach argument quality as a point-wise ranking task, with the goal of selecting argument segments that clearly express an argument facet in a given dialogue. Arguments are labeled by a real value in the range of  $[0, 1]$ , where a score of 1 indicates that an argument can be easily interpreted. They then develop an automatic regression method using these labels. Their corpus, which we refer to henceforth as *SwanRank*, contains 5.3k labeled arguments.

An alternative approach to assess arguments is to focus on their relative *convincingness*, by comparing pairs of arguments with similar stance. This approach is introduced in ? (?; ?), and further assessed in ? (?), ? (?), and ? (?). As part of their work, ? (?) introduce two datasets: *UKPConvArgRank* (henceforth, *UKPRank*) and *UKPConvArgAll*, which contain 1k and 16k arguments and argument-pairs, respectively. In their work, the point-wise scores are induced from the pair-wise labels, rather than annotated directly. ? (?) focus of ranking convincingness of *evidence*. Their solution is based on a Siamese neural network, which outperforms

the results achieved in ? (?) on the *UKP* datasets, as well as several baselines on their own dataset, *IBM-ConvEnv*. ? (?) present a method that is based on representing an argument by the sum of its token embeddings, extended in ? (?) to include a Feed Forward Neural Network. These two works outperform ? (?) on the *UKP* datasets for both the pair-wise and point-wise tasks, respectively.

? (?) present a new dataset comprised of over 47k claims in 471 topics from the website *kialo.com*, aimed at evaluating the effect of pragmatic and discourse context when determining argument quality. They propose models to predict the *impact* value of each claim, as determined by the users of the website. Their dataset is somewhat different from ours as it focuses on argument impact, rather than overall quality, and doing so in the context of an argumentative structure, instead of independently. In addition, their impact values are based on spontaneous input from users of the website, whereas our dataset was carefully annotated with clear guidelines. Still, it further highlights the importance of this field.<sup>2</sup>

? (?) consider both point-wise and pair-wise quality approaches, as well as the interaction between them. They introduce two datasets: *IBMRank*, which contains 5.3k point-wise labeled arguments (6.3k before cleansing) and *IBM-Pairs*, which contains 9.1k labeled argument-pairs (14k before cleansing). Arguments in their datasets, collected in the context of Speech by Crowd experiments, are suited to the use-case of civic engagement platforms, giving premium to the usability of an argument in oral communication. Our dataset differs in three respects: (1) our dataset is larger by a factor of 5 compared to previous datasets annotated for point-wise quality; (2) our data were collected mainly from crowd contributors that presumably better represent the general population compared to targeted audiences such as debate clubs; (3) we performed an extensive analysis of argument scoring methods and introduce superior scoring methods that consider annotators credibility without removing them entirely from the labeled data, as is done in ? (?).

## 3 IBM-Rank-30k Dataset

In the next two sections, we present the creation of the *IBM-Rank-30k* dataset. First, we describe the process of argument collection. We then move to describing how arguments were annotated for quality. Finally, we discuss and analyze how we derive point-wise continuous quality labels from binary annotations, by conducting a comprehensive comparison between different scoring functions. We release this dataset as part of this work.<sup>3</sup>

### 3.1 Argument Collection

For the purpose of collecting arguments for the *IBM-Rank-30k* dataset, we conducted a crowd annotation task, using the Figure Eight platform.<sup>4</sup> A small portion of arguments (8.6%) was also collected from expert annotators who work closely

<sup>2</sup>The work by ? (?) was published after our submission, therefore we were not able to fully address it.

<sup>3</sup><http://ibm.biz/debater-datasets>

<sup>4</sup><http://figure-eight.com/>

with our team. We selected 71 common controversial topics, for which arguments were collected (e.g., *We should abolish capital punishment*).

We follow similar guidelines to that presented in ? (?). Annotators were presented with a single topic each time, and asked to contribute one supporting and one contesting argument for it, requiring arguments to be written using original language. To motivate high-quality contributions, contributors were informed they will receive extra payment for high quality arguments, as determined by the subsequent argument quality labeling task (Section 3.2). It was explained that an argument will be considered as a high-quality one, if a person preparing a speech on the topic will be likely to use this argument *as is* in her speech.

Similarly to ? (?), we place a limit on argument length - a minimum of 35 characters and a maximum of 210 characters. In total, we collected 30,497 arguments from 280 contributors, each contributing no more than 6 arguments per topic.

### 3.2 Argument Quality Annotations Collection

In this section we describe the argument quality annotation process, performed for all collected arguments. As above, we used the Figure Eight platform, with 10 annotators per argument. Following ? (?), annotators were presented with a binary question per argument, asking if they would recommend a friend to use that argument *as is* in a speech supporting/contesting the topic, regardless of personal opinion. In addition, annotators were asked to mark the stance of the argument towards the topic (*pro* or *con*).

To monitor and ensure the quality of the collected annotations, we employed the following measures introduced in ? (?):

**Test Questions.** Before the labeling of 1/5 of the arguments, a hidden test question about the stance of the argument towards the topic was presented, aimed to verify the annotator is reading the argument carefully. Annotators that failed more than 20% of the test questions were removed from the task, and their judgments were ignored. Typically 5% – 10% of the contributors were removed from each sub-task due to this reason.

**Annotator-reliability score.** Defined in ? (?) as the Annotator- $\kappa$  for a single score (and task-average- $\kappa$  averaged on all valid scores) and denoted here as *Annotator-Rel*. This score was used both to monitor tasks in real time, and as a basis for the weighted score function described in Section 4. It is obtained by averaging all pair-wise  $\kappa$  for a given annotator, with other annotators that share at least 50 common judgements. Annotators who do not share at least 50 common judgments with at least 5 other annotators, do not receive a value for this score.

The average of valid annotators reliability scores on the quality annotations is 0.12. This task reproduces the task from ? (?), and as established there, such an average is acceptable due to the subjectivity of the task.<sup>5</sup> Among other things, ? (?) rely on the task-average- $\kappa$  of the stance annotations in the same task, which was 0.69 in that work, and

here it is 0.83.

To further ensure high quality annotations, rather than introducing the task to any crowd worker, we used a selected group of 600 crowd annotators, which had high Annotator-reliability in past tasks of our team.

## 4 Deriving an Argument Quality Score from Binary Annotations

As mentioned in section 3.2, for each argument several annotators answered a binary question regarding its quality. We chose this format to simplify the annotation process of this subjective question, aiming to avoid an additional subjective element - the scale. However, we still need to provide a single continuous score per argument, reflecting its quality, that can be compared with the scores of other arguments. To that end, we evaluate approaches to derive such a quality score, on a bounded scale, from a set of binary annotations.

### 4.1 Quality Scoring Functions

First, we describe two scoring functions. Each function provides the likelihood of the positive label, between 0 (bad quality) and 1 (good quality).

**MACE probability (MACE-P)** - ? (?) suggested MACE (?) as a scoring function for the quality of an argument based on crowd annotations. MACE is an unsupervised item-response generative model which predicts the probability for each label given the annotations. MACE also estimates a reliability score for each annotator which it then uses to weigh this annotator’s judgments. We use the probability MACE outputs for the positive label as the *MACE-P* scoring function.

**Weighted-Average (WA)** - We suggest to use a weighted-average score to incorporate annotator-reliability, in the spirit of MACE. This is designed to decrease the influence of non-reliable annotators on the final quality score, thus providing an intuitive and gradual form of data cleansing. For each argument  $a$ , we define  $P_a$  as the set of annotators who labeled it as positive, and  $N_a$  as the set of annotators who labeled it as negative. The WA score of an argument  $a$  is defined as:

$$Score(a) = \frac{\sum_{Annotator_i \in P_a} Annotator-Rel_i}{\sum_{Annotator_j \in N_a + P_a} Annotator-Rel_j}$$

There is a clear distinction between the distribution of scores obtained by WA and *MACE-P* scoring functions (see Figure 1). WA outputs values close to 0 or 1 only if there is a strong annotation consensus. As generally there are more positive annotations in our data, the histogram of this scoring function is skewed towards 1, with an almost linear decrease. On the other hand, MACE assigns probabilities to both labels. As a result, the quality scores lean strongly to both extreme values, creating a U-shaped histogram.

### 4.2 Comparing Scoring Functions

Next, we compare the quality scoring functions described above via the following three experiments.

<sup>5</sup>In ? (?) it was 0.1.




Figure 1: Histogram of arguments according to *MACE-P* and *WA* quality scores. X axis - quality scores. Y axis - counts of arguments in *IBM-Rank-30k*.

**Disagreements in choosing the better argument** We first ask which of the two methods is ‘correct’ more often when examining the cases in which given a pair of arguments, they disagree on which argument should get a higher score. We created all possible argument pairs from *IBM-Rank-30k*, such that arguments in the same pair are taken from the same topic and have the same stance towards it (to avoid annotator bias).<sup>6</sup>

We focus on the set of pairs in which the two methods disagree on the preferred argument, i.e. the one that received a higher score (about 20% of the pairs). We sample 850 such disagreement pairs and send them for *pair-wise* annotation, i.e., asking the annotators to choose the better of the two arguments. Each pair was annotated by 7 annotators from our selected group of crowd annotators. We discarded pairs for which the agreement between annotators was less than 70% (this filtered out 27% of the pairs). In 55% of the remaining pairs the annotators chose the argument preferred by *MACE-P*, making this method somewhat more correlated with the pair-wise judgments. Both methods were also compared to simple-average, the scoring method applied in ? (?). In 61% of the pairs that differ between simple-average and *MACE-P*, annotators chose the argument preferred by the latter. In 59% of the pairs that differ between simple-average and *WA*, annotators chose the argument preferred by *WA*. As the simple-average method was found inferior, it was omitted from subsequent evaluation experiments.

**Agreement with pair-wise annotation** In this experiment, we present each scoring function with pairs of arguments, and compare their preferred arguments with a gold standard obtained by pair-wise annotation, following the consistency evaluation performed in ? (?). Similar to the previous experiment, we generate all argument pairs (of the same topic and stance). Next, we bin all pairs into four sets by the size of the delta between the scores of their arguments

<sup>6</sup>Arguments were selected from the training set, which is detailed in Section 6.1.

(e.g., all pairs with score difference between 0.25 and 0.5). From each such *delta bin* we randomly sample 150 pairs and send them for pair-wise annotation as described in the previous experiment. We repeat this process for both scoring functions.

We calculate the precision of each scoring function in each delta bin, based on its agreement with the pair-wise annotation. Table 2 depicts the results for *MACE-P*, and Table 3 for *WA*. The results show that as the score difference increases (larger *delta bin*), precision also increases (for both scoring functions). Interestingly, this tendency is more prominent for *WA*, reaching a perfect match when the difference in point-wise quality scores is higher than 0.75.<sup>7</sup> The tables also depict the percent of pairs filtered due to low (less than 70%) agreement between annotators. Interestingly, in the 3 bins with more than 0.25 difference in point-wise scores, less pairs are filtered when using *WA*, suggesting that when *WA* is confident about which argument is better, annotators also tend to agree more often, which is not the case for *MACE-P*.

Delta bin	Filtered pairs	Precision
$\leq 0.25$	15%	0.64
$(0.25, 0.50]$	23%	0.77
$(0.50, 0.75]$	14%	0.81
$> 0.75$	13%	0.88

Table 2: Comparing *MACE-P* scoring function preference to gold standard of pair-wise annotation.

Delta bin	Filtered pairs	Precision
$\leq 0.25$	15%	0.66
$(0.25, 0.50]$	13%	0.74
$(0.50, 0.75]$	6%	0.90
$> 0.75$	3%	1.00

Table 3: Comparing *WA* scoring function preference to gold standard of pair-wise annotation.

**Split annotations consistency** A desirable property for a scoring function is that it will be relatively consistent with respect to different sets of annotators; namely, if we split the binary annotations into two sets and construct the continuous annotation from each set independently, we will end up with approximately the same score.

To examine that, we randomly split the 10 binary annotations we have for each argument into two sets (each containing 5 annotations), similar to ? (?). We then calculate two quality scores for each argument, based on these two smaller annotation sets, respectively. While ? (?) utilized agreement to measure consistency, we can measure correlation between the two continuous scores.

We note that since each score was calculated on a set which is half the size of the original annotation set, we have

<sup>7</sup>However, we note that while for *MACE-P* the  $> .75$  delta bin holds for 19.5% of the pairs, for *WA* it holds less than 1%.

less information for assessing the reliability score of each annotator. This fact can harm the accuracy of the scoring functions, as they both utilize annotators reliability scores. Nevertheless, for both scoring functions we considered, we find a good correlation between the scores calculated over the two smaller annotation sets. *WA* achieves 0.42 and 0.36 Pearson and Spearman correlations, respectively. *MACE-P* achieves 0.42 in both.

In summary, we point out a key difference between the two scoring functions: the tendency of *WA* to present a gradual continuous scale, as opposed to *MACE-P*, which aims at discovering the ‘true’, hence the binary, labels. For this reason we tend to prefer *WA* as a scoring function for this task, which is inherently deriving a non-binary score. However, as our experiments do not show a clear preferred function, we utilize both for the evaluation of neural methods in Section 6.1. For brevity, in the analyses in sections 5 and 7 we only use *WA* scores. Finally, in the dataset we release as part of this work, we include the quality scores of both scoring functions. We conclude that the induction of a quality score on top of existing annotations is not necessarily trivial, and should be carefully considered, as it impacts the scores distribution and the performance of learning algorithms trained on the score.

## 5 Analysis of Quality Dimensions

We seek to further explore the consistency and accuracy of the *IBM-Rank-30k* dataset. To that end, we employ a quality dimensions model, i.e. a model that decomposes a holistic quality score to several dimensions. Such a model was suggested by ? (?) - a meta-model that was created via a broad literature survey on previous quality models. They decompose quality to 15 sub-dimensions that aim to determine fine-grained properties of argument quality.

We conducted an annotation task similar in spirit to ? (?), in the context of the *IBM-Rank-30k* dataset. Our goal is to quantitatively assess the reasons for arguments having low or high quality, through the prism of this theoretical model. First, we decided to exclude 5 dimensions from this task, as they lack high potential to embody relevant characteristics over our data.<sup>8</sup> We split the *IBM-Rank-30k* dataset to 5 equally populated bins according to the *WA* quality scores (1-5, where 1 is the lowest quality bin), and randomly sample 100 arguments with a uniform distribution over the bins. Each argument was labeled by 3 expert annotators that have extensive background in related tasks of our team. The annotators were not aware of the original quality bins. We have asked the annotators to annotate each argument according to each of the 10 dimensions on a scale of 1-3, to stay consistent with the scale offered by ? (?), and calculate the average for each dimension.

We observe that even though the task is complex, across all dimensions, the arguments from the highest bin achieve a higher average over the middle bins, and arguments from the middle bins achieve a higher score over the lower bins. The 2 dimensions that present the largest difference between bins

<sup>8</sup>Those 5 dimensions are: Appropriateness, Arrangement, Credibility, Emotional Appeal and Reasonableness.

5 and 1 are *Global Relevance* and *Effectiveness*, with an average difference of 0.72 and 0.64, respectively. *Global Relevance* asks whether an argument provides information that helps to arrive at an ultimate conclusion regarding the discussed issue, while *Effectiveness* asks whether the argument is effective in helping to persuade in the author’s stance. These results suggest that the differences between low and high arguments in the *IBM-Rank-30k* dataset are best explained by how related the annotator found the argument to the topic, and how effectively the argument was presented. Such results corroborate several notions. Firstly, our quality labeling is broadly consistent with a known quality model. Moreover, the outcome serves as an initial proof of concept of decomposing quality in a large dataset to gain explainability. The detection of quality dimensions can be used in a range of applications, from more exact research on what determines quality, to feedback systems that can recommend users more precisely what they need to improve to advance their argumentation skills.

## 6 Argument Quality Ranking

We now move to apply the *IBM-Rank-30k* dataset to the task of learning to rank the quality of arguments. For this purpose, we evaluate the following methods, which include several neural methods, as well as some simpler baselines.

**Arg-Length.** Although we placed a strong limit on argument length, it is possible that there is still a bias towards longer arguments. To inspect this, we evaluate a ranking baseline based on an argument’s length in characters.

**SVR BOW.** We evaluate a Support Vector Regression ranker, implemented by the scikit-learn toolkit,<sup>9</sup> with an RBF kernel and bag-of-words features, using the most frequent 1000 tokens in the training set.

**Bi-LSTM GloVe.** As a simple neural baseline we implement a Bi-LSTM model with self-attention, following the model used in ? (?). The model was trained with a dropout of 0.15, an LSTM layer of size 128 and an attention layer of size 100. For input features we used the 300 dimensional GloVe embeddings (?).

We use the following three methods based on BERT:

**BERT-Vanilla** (henceforth, **BERT-V**). In this network, for each argument, we concatenate the last 4 layers of the [CLS] token obtained from BERT’s pre-trained model, resulting in a feature vector of size  $4 \times 768 = 3072$ . The features are passed through a fully-connected hidden layer of size 100 with ReLU activation, after which we apply a sigmoid activation layer with a single output.

**BERT-Finetune** (henceforth, **BERT-FT**). This method fine-tunes BERT’s pre-trained model. The official code repository of BERT<sup>10</sup> supports fine-tuning to classification tasks, which is done by applying a linear layer on the [CLS] token of the last layer of BERT’s model, which is then passed through a soft-max layer. The weights of the preceding layers are initialized with BERT’s pre-trained model, and the entire network is then trained on the new data. To

<sup>9</sup><https://scikit-learn.org/>

<sup>10</sup><https://github.com/google-research/bert>

adapt the fine-tuning process to a regression task, the following were performed: (1) Changing the label type to represent real values instead of integers; (2) Replacing the softmax layer with a sigmoid function, to support a single output holding values in the range of  $[0, 1]$ ; (3) Modifying the loss function to calculate the Mean Squared Error of the logits compared to the labels.<sup>11</sup>

**BERT-FT<sub>TOPIC</sub>.** We also evaluate the addition of the topic to the input of BERT-FT. The topic is concatenated to the argument, separated by a `[SEP]` delimiter, and the model is fine-tuned as in BERT-FT.

### 6.1 Experiments on *IBM-Rank-30k*

For the purpose of evaluating our methods on the *IBM-Rank-30k* dataset, we split its 71 topics to 49 topics for training, 7 for tuning hyper-parameters and determining early stopping (dev set) and 15 for test.

We present results for both *WA* and *MACE-P* scoring functions, aiming to shed some more light on their properties. All models were trained for 5 epochs over the training data, taking the best checkpoint according to the performance on the dev set, with a batch size of 32 and a learning rate of  $2e-5$ . We calculate Pearson ( $r$ ) and Spearman ( $\rho$ ) correlations on the entire test set.

For significance testing, we use the Williams test (?) which evaluates the significance of a difference in dependent correlations (?). The Williams test has been successfully used in Machine Translation in the evaluation of MT metrics (?) and quality estimation (?).

### 6.2 Results and Discussion

The results on the *IBM-Rank-30k* dataset are presented in Table 4. Using BERT-V improves on the Bi-LSTM GloVe method by .4-.6 points for Pearson correlation, and .2-.6 points for Spearman correlation. Fine-tuning BERT improves on BERT-V by .2-.4 points for both correlation measures. Adding the topic adds a statistically significant improvement of .1-.2 points ( $p \ll 0.01$  for both correlation measures and quality score methods). Interestingly, when using *MACE-P* scores, the model is able to achieve higher Spearman correlation for most methods. Also, argument length is not an indicator for quality, as evident by the poor performance of the Arg-Length baseline.

	<i>WA</i>		<i>MACE-P</i>	
	$r$	$\rho$	$r$	$\rho$
Arg-Length	.21	.22	.22	.23
SVR BOW	.32	.31	.33	.33
Bi-LSTM GloVe	.44	.41	.43	.42
BERT-V	.48	.43	.49	.48
BERT-FT	.51	.47	.52	.5
BERT-FT <sub>TOPIC</sub>	<b>.52</b>	<b>.48</b>	<b>.53</b>	<b>.52</b>

Table 4: Correlations on the *IBM-Rank-30k* test set.

<sup>11</sup>The implementation of BERT fine-tuning in regression mode is adapted from: <https://github.com/google-research/bert/issues/160#issuecomment-445066341>.

An important property of a good model is that its performance increases when considering arguments that are on the extremes of the argument quality scale. To evaluate this, we define a *cut-off* percentile  $d$  as a view of the test data that considers only the top  $d$  and bottom  $d$  percent of the data. For each cut-off we calculate Pearson and Spearman correlations w.r.t the predictions of the BERT-FT<sub>TOPIC</sub> model. Figure 2 presents the correlations for cut-off percentiles ranging from 10% to 50% (50% is equivalent to taking the entire data), for the models trained on *MACE-P* and *WA* quality scores. A clear trend emerges in which the correlations increase as arguments are taken from a smaller percentile, i.e. from further extremes of the quality scale, reaching up to .71-.73 and .67 for Pearson and Spearman correlations, respectively, when considering only the bottom and top 10% of the test set for evaluation.




Figure 2: Pearson (left) and Spearman correlations of various cut-off percentiles for the BERT-FT<sub>TOPIC</sub> models, trained on data containing *WA* and *MACE-P* quality scores.

### 6.3 Experiments on *SwanRank* and *UKPRank*

In this section we demonstrate the strength of the two best methods – BERT-FT and BERT-FT<sub>TOPIC</sub> – by applying them on two related datasets, *SwanRank* and *UKPRank*. We chose *SwanRank* as it also contains direct point-wise quality labels. *UKPRank* was chosen as an established dataset which also contains point-wise quality scores. However, it should be noted that as opposed to the *IBM-Rank-30k* and *SwanRank* datasets, the point-wise labels in *UKPRank* were induced from a pair-wise labeling task, and not obtained directly, as mentioned previously.

**SwanRank dataset.** In ? (?) the *SwanRank* dataset was evaluated both in in-domain and cross-domain scenarios. In this work, we focus on the latter, more challenging scenario. Their model is a SVR with RBF kernel, with features coming from one of many feature sets they experiment with. They randomly split each topic to train (75%) and test sets, and run their model, with each possible feature set, on each train-test cross-domain pair. That is, the topics of the train and the test sets are distinct. We adapt to their approach as follows:

- For each topic, we create our own random split to train and test, as the exact split to train and test for each topic was not available to us.
- For each topic, we consider its best result achieved in ? (?) in the cross-domain scenario. For example, for evaluating the *gun control* topic, the best result in ? (?) is obtained by training on the *gay marriage* topic. By this we obtain 4 train-test pairs: *gay-marriage* (GM, train)–*gun control* (GC, test), *gun-control*–*gay marriage*, *death penalty* (DP)–*evolution* (EV), and *evolution*–*death penalty*. It should be noted, that adapting to this setup puts our work at a disadvantage, because the best results for different pairs in ? (?) are achieved by using different feature sets, as opposed to a single learning framework. Thus, this setting represents an upper limit rather than an actual obtained result.
- For each of the 4 pairs, we train our BERT-FT and BERT-FT<sub>TOPIC</sub> methods for 2 epochs on the train topic, and test the model on the test topic. We compute Root Relative Squared Error (rrse), following ? (?), and report results for each pair as well as the weighted-average.

**Results.** Results on the *SwanRank* dataset are presented in Table 5. Using the BERT-FT<sub>TOPIC</sub> method yielded an average improvement of .8 points compared to the optimal result taken from ? (?). Performance has improved in 3/4 test topics (*gay marriage*, *gun control* and *death penalty*), and decreased in the *evolution* topic. However, in ? (?) the performance on the *death penalty* and *evolution* topics is low even in the easier in-domain task, presumably indicating they are much more difficult to predict.

Train Test	GC GM	GM GC	EV DP	DP EV	Avg
Swanson	0.84	0.81	1	<b>0.97</b>	0.89
BERT-FT	0.8	0.82	1.1	0.98	0.91
BERT-FT <sub>TOPIC</sub>	<b>0.65</b>	<b>0.71</b>	<b>0.96</b>	1.01	<b>0.81</b>

Table 5: Weighted-average RRSE on the 4 topic pairs of the *SwanRank* dataset (a lower score is better). *Swanson* row: the result by averaging the best train-test pairs cross-domain results published in ? (?).

**UKPRank dataset.** We conduct cross-validation where in each fold we trained on 31 topics and tested on the held-out topic. The model was evaluated after 5 training epochs. Following ? (?), we report average Pearson ( $r$ ) and Spearman ( $\rho$ ) correlations, and compare the results of our methods to the Bi-LSTM and GPPL methods published there, as well as to the EviConvNet method, the best result from ? (?), and to the SWE+FFNN method, the best result from ? (?).

**Results.** The results on the *UKPRank* dataset are presented in Table 6. Both of our methods obtain Pearson correlations which are comparable to the SWE+FFNN, EviConvNet, and GPPL methods, and worse Spearman correlations. However, as the point-wise quality scores were obtained via a pair-wise proxy, previous methods assume the existence of a pair-wise labeled dataset for training, and EviConvNet, for example, trains on the pair-wise labels directly. Our method

is less suitable for this setting, as it does not depend on any pairs being labeled for training, making this comparison less trivial.

	$r$	$\rho$
Bi-LSTM GloVe	.32	.37
Bi-LSTM ling+GloVe	.37	.43
GPPL ling	.38	.62
GPPL GloVe	.33	.44
GPPL ling+GloVe	.45	.65
GPPL opt. ling+GloVe	.44	.67
EviConvNet	.47	.67
SWE+FFNN	<b>.48</b>	<b>.69</b>
BERT-FT	.45	.63
BERT-FT <sub>TOPIC</sub>	.46	.62

Table 6: Average correlation on the *UKPRank* dataset.

## 7 Learning to Represent Quality

During fine-tuning BERT to a new task, the weights of the pre-trained model are updated, and as a result, the contextual representations of tokens change. In this section, we exemplify how these new embeddings are enriched with contextual quality properties. Using an anecdotal example, we show how quality is encoded in BERT’s token embeddings, after being exposed to the *IBM-Rank-30k* dataset during training. We do this by comparing the contextual embeddings retrieved from BERT’s pre-trained model and from the BERT-FT model, of a common token in the *IBM-Rank-30k* dataset, *people*. For the purpose of this analysis, we use the model fine-tuned for the *IBM-Rank-30k* dataset with *WA* scores. First, we split the data to 5 equally populated bins, according to the *WA* quality score.<sup>12</sup> We then retrieve the contextual embeddings of *people* in a sample of 20 arguments which contain it, 10 from the lowest and 10 from the highest quality bins. Embeddings are taken from the last layer of BERT’s model before and after fine-tuning, resulting in two sets of 20 embeddings overall. We cluster each set of embeddings using Hartigan’s K-Means (?; ?), and calculate the adjusted mutual information (AMI) of the clusters with respect to the low and high quality bins. When clustering the embeddings retrieved from BERT’s pre-trained model, the AMI is 0.08, a low result which is expected given that BERT’s language model should not have any preference to argument quality. However, when clustering the embeddings extracted from BERT-FT, the AMI reaches 0.31, indicating that these representations absorb the qualitative nature of the arguments they come from. We use t-SNE (?) to visualize these embeddings in a 2D space (Figure 3). Before fine-tuning, the embeddings are clustered together, whereas after fine-tuning, the embeddings are partially separated by the quality of the arguments, further indicating that the quality contributes to the contextual representation of this token.

<sup>12</sup>We exclude the training set, to avoid a trivial bias of the BERT-FT model having trained on this set.




Figure 3: A 2D t-SNE projection of the embeddings of the term `people` taken from a sample of high (x) and low (o) quality arguments from the dev and test sets of the *IBM-Rank-30k* dataset. Left - using BERT’s pre-trained model. Right - after fine-tuning.

## 8 Conclusions and Future Work

In the rapidly expanding field of computational argumentation, argument quality is an increasingly prominent issue, having substantial implications. To advance the development of argument quality ranking models, the creation of new datasets is a critical step. In this work, we present a novel argument dataset labeled for point-wise quality, *IBM-Rank-30k*, containing 30,497 arguments. To the best of our knowledge, this dataset is the largest to include point-wise quality labels, 5 times larger than previously released datasets. We follow ? (?) by collecting the arguments actively, while employing elaborate annotation control measures. A practical question, overlooked in previous datasets, is how to induce continuous labels from binary annotations. We address this issue by conducting an extensive comparison of two common approaches and analyzing their appropriateness to our dataset. We also exploit this dataset to the task of argument quality ranking, by presenting a BERT-based neural method which outperforms several baselines. We show this method is capable of achieving promising results on other datasets as well. We believe that the approach to argument collection, the analysis of different labeling scores, as well as the sheer size of the dataset, make it useful for further advancements in this field.

As an attempt to provide insight regarding the characteristics of *IBM-Rank-30k*, we conducted an analysis of quality dimensions, showing that the dimensions of *Global Relevance* and *Effectiveness* are the most indicative to overall quality scores. As future work, we would like to further explore this, by investigating how quality dimensions impact overall quality, and whether a prediction model can capture these dimensions effectively.

## Acknowledgements

We thank Eyal Shnarch, Leshem Choshen, and Elad Venezian for their insightful comments and suggestions.

Eum dolore rerum et, reiciendis eveniet ullam aut illum inventore odio earum aliquam. Itaque non incidunt iusto aut, molestiae quos ipsa alias quae debitis sit repudiandae, sint tempora totam dolore ipsa, dolorum aspernatur accusamus minima voluptatibus numquam nisi porro exercitationem ipsum facere? Sapiente eveniet pariatur ea, fugiat ipsa vel recusandae impedit hic commodi ab dignissimos soluta, odio quibusdam vero dolores, tempora libero quam facilis placeat doloremque accusamus expedita autem deserunt omnis sint? Provident repellendus laudantium, voluptates totam

numquam iure error molestiae voluptas, repellat molestiae veritatis dolorum, minima rerum dolores animi harum tempora modi quis, voluptatem laudantium necessitatibus minus nulla. Dolor quam qui voluptate, minus rerum quos ducimus praesentium aliquid delectus assumenda nulla, cumque exercitationem voluptates fugiat dolorem, possimus accusantium aspernatur qui. Repellendus maiores exercitationem aliquid vel molestiae reiciendis labore quidem laudantium dicta, fugiat mollitia neque ipsum officiis hic molestiae eligendi recusandae, enim debitis expedita quasi non suscipit tenetur officiis, hic sapiente rerum sequi voluptatibus doloribus temporibus dicta illum similique quis quam, saepe nulla quae dicta sequi deserunt doloremque officia quidem provident dolorem? Fugiat sunt nesciunt eaque voluptatibus porro dignissimos, illo ducimus omnis ipsa totam, saepe natus nisi quo quaerat reiciendis nemo placeat ab perferendis numquam, optio velit possimus provident in laboriosam distinctio culpa reiciendis totam alias, ducimus voluptatem laborum nesciunt assumenda dolorum veniam? Quibusdam aperiam minima vero dolore consequuntur, mollitia voluptate molestias harum aliquid ullam suscipit voluptatem, ipsum autem doloremque iste velit distinctio? Explicabo deleniti nostrum necessitatibus illo dolor architecto velit voluptatibus, quia officiis eius unde enim nemo quos fuga eaque alias, voluptatem maiores ipsa facilis consequuntur. Esse fugit blanditiis ab minima aspernatur vel ex nulla inventore ipsa, delectus reprehenderit minus corrupti nesciunt inventore possimus quasi? Dicta perspiciatis vel nihil nulla aspernatur explicabo soluta obcaecati quibusdam provident, repudiandae veritatis tempore consequuntur, dignissimos praesentium molestias veritatis eius voluptate magni perspiciatis ipsa corporis perferendis hic, excepturi eveniet sapiente perspiciatis cumque eos neque porro ad molestias, ex quo officia eveniet inventore voluptate. Consectetur nobis aliquid quisquam doloribus sapiente saepe veritatis explicabo nam, non molestiae ea nulla voluptas vitae porro dolores, totam repellendus dicta quas molestias quo ut eveniet ex, totam ipsam voluptatibus magni neque enim molestiae eum sed? Eligendi aspernatur id iure corporis possimus minus, voluptatum dignissimos in aut necessitatibus vel amet minima provident mollitia obcaecati exercitationem, eaque ea blanditiis aut odit porro molestiae maiores, qui voluptate veniam totam minus doloremque debitis ratione reiciendis ut, itaque magnam aspernatur repellat fugiat magni doloremque impedit? Quas eos quidem quia corporis, incidunt quam tempore minus nisi corrupti nostrum pariatur dolores repudiandae, alias vero quod quas reiciendis a fuga sequi, et cum iusto quam minima architecto quos mollitia reiciendis. Ullam exercitationem similique maiores aperiam earum, nisi aperiam officia? Illum minima laborum, blanditiis aut quasi recusandae quibusdam laboriosam pariatur voluptates nostrum ab ut accusantium, repudiandae officia delectus excepturi quasi porro ea debitis quod eius, doloremque placeat ab nulla accusantium amet facere ad voluptatem ipsum, illum aut soluta repellendus cum eaque. Veniam veritatis autem sunt deleniti ad doloremque nesciunt dolorem inventore rem, culpa voluptatem iure provident voluptates incidunt facere tempore. Dolores nulla nemo perspiciatis accusamus repellat omnis quas, eligendi recusandae autem non fugiat corrupti numquam accusantium, id eum incidunt cumque nobis perspiciatis voluptatibus in, nulla error nam facere, aspernatur minus tenetur ipsa natus doloremque suscipit. Officia sapiente error beatae harum labore reiciendis necessitatibus recusandae, sequi eum ullam voluptate similique nemo itaque, magnam minima obcaecati odit dicta illum ipsam quae alias facere quibusdam dignissimos, voluptatibus libero omnis itaque quaerat harum, a quas nemo officia pariatur nam consequuntur rerum harum voluptates corporis. Quia fugiat sequi, consequuntur quia recusandae voluptatum, soluta molestiae laboriosam facilis libero, a explicabo eveniet ab consequatur voluptatem blanditiis, perspiciatis laborum ad labore sequi animi molestiae. Est delectus eos id similique, nam eos



at obcaecati, dolores minima eaque harum ipsum labore ipsam, est corrupti odit quam.Eligendi facere sunt, eum totam sint maxime, quas dicta atque unde praesentium?