

Complexity of Probabilistic Inference in Random Dichotomous Hedonic Games

Saar Cohen, Noa Agmon

Department of Computer Science, Bar-Ilan University, Israel
saar30@gmail.com, agmon@cs.biu.ac.il

Abstract

Hedonic games model cooperative games where agents desire to form coalitions, and only care about the composition of the coalitions of which they are members. Focusing on various classes of dichotomous hedonic games, where each agent either approves or disapproves a given coalition, we propose the random extension, where players have an independent participation probability. We initiate the research on the computational complexity of computing the probability that coalitions and partitions are optimal or stable. While some cases admit efficient algorithms (e.g., agents approve only few coalitions), they become computationally hard ($\#P$ -hard) in their complementary scenario. We then investigate the distribution of coalitions in perfect partitions and their performance in majority games, where an agent approves coalitions in which the agent is friends with the majority of its members. When friendships independently form with a constant probability, we prove that the number of coalitions of size 3 converges in distribution to a Poisson random variable.

Introduction

In several real-life scenarios arising from economics, politics, and sociology, we notice the phenomenon of coalition formation, where each person, termed as agent, forms coalitions with others to get some benefit, experiencing a utility that depends on the particular set of agents she joins. A popular game-theoretic approach to the study of coalition formation problems is Hedonic games (Dreze and Greenberg 1980), whose outcome is a partition of the agents into coalitions, over which the agents have preferences. One of their main properties is non-externality: an agent minds only her own coalition, regardless of how the others aggregate. Yet, the number of coalitions an agent can be part of is exponential in the number of agents, and therefore it is desirable to consider expressive, but succinctly representable classes of hedonic games. One common approach is to restrict the players' possible preferences. As such, our work focuses on dichotomous preferences, where each agent ei-

ther approves or disapproves of a given coalition (Aziz et al. 2016).

Two major underlying assumptions in hedonic games literature are that all agents are assured to participate in the game, and the nature of their collaborations and preferences over those are certain. However, both often fail to hold for real-world problems. Generally, participation in the game does not depend only on strategic choices, but also on external factors. For instance, in a sports tournament, players may not arrive to the game before its inception due to a last minute injury, weather conditions or traffic jams.

In this paper, we thus lay the theoretical foundations for studying the interplay between dichotomous hedonic games (DHGs) and the complexity of probabilistic inference. Motivated by probabilistic inference in voting games (Bachrach et al. 2011; Imber and Kimelfeld 2021), we propose random DHGs (RDHGs), where each player has an independent participation probability. For various classes of DHGs, we explore the computational complexity of computing the probability that coalitions and partitions are either optimal, perfect or stable. We first examine coalitions and partitions that maximize the number of approving players within them. Then, we regard partitions which are either stable or perfect (where every agent is in an approved coalition). We present cases that admit poly-time algorithms for some of these problems: when agents approve at most one coalition; when agents approve at most one coalition size; when agents are placed on a line and only approve intervals. Yet, the above problems become computationally hard ($\#P$ -hard) when the setting becomes non-linear. Accordingly, we show cases where deciding whether the probability of optimality is nonzero can be computed in poly-time, and computing the probability of nonoptimality can be done up to a multiplicative approximation. Though DHGs are one of the very few sub-classes of hedonic games that admit poly-time solutions (Peters 2016), our results put this property of DHGs in question when it comes to a probabilistic setting.

Our model corresponds to counting variants of a new type of manipulation in hedonic games: constructive control by adding players. Analogous to election con-

Class	Polynomial	#P-Hard	NP-Complete
1-Lists	$\mathbf{C} \subseteq N$ is $\text{WO}_{\mathcal{C}}$; π is WO_{Π}	WO_{Π} ; <u>SCS</u>	π is $\text{WO}_{\Pi} > \mathbf{0}$; $\text{SCS} > \mathbf{0}$
k -Lists	$\mathbf{C} \subseteq N$ is $\text{WO}_{\mathcal{C}} > \mathbf{0}$; π is $\text{WO}_{\Pi} > \mathbf{0}$	$\mathbf{C} \subseteq N$ is $\text{WO}_{\mathcal{C}}$; WO_{Π} ; <u>SCS</u>	$\text{WO}_{\Pi} > \mathbf{0}$; π is $\text{WO}_{\Pi} > \mathbf{0}$; $\text{SCS} > \mathbf{0}$
AHG	π is PF; π is NS	<u>PF</u> ; <u>NS</u> ; <u>SCS</u>	$\text{PF} > \mathbf{0}$; $\text{NS} > \mathbf{0}$; $\text{SCS} > \mathbf{0}$
CI	$\mathbf{C} \subseteq N$ is $\text{WO}_{\mathcal{C}}$; π is WO_{Π} ; π is PF		
RG			$\text{PF} > \mathbf{0}$; $\text{NS} > \mathbf{0}$

Table 1: Overview of complexity results for various dichotomous preferences. We consider existence of welfare-optimal coalitions or partitions ($\text{WO}_{\mathcal{C}}$ and WO_{Π} , resp.), and perfect (PF), Nash-stable (NS) and strict-core-stable (SCS) partitions. Results in boldface apply to both probabilistic inference and the induced control problem, underlined ones only apply to the control problem and the remaining only apply to probabilistic inference. Considering the "Preliminaries" section, the counting problems (Problem 1) relate to probability computation, whereas decision problems (Problem 3) correspond to verifying its zeroness.

trol problems (Bartholdi et al. 1992), its goal is ensuring that a specified outcome satisfies a certain solution concept by adding players. Unlike Sybil attacks (Vallée et al. 2014), players are sincere and participate probabilistically. When players approve polynomially many coalitions, we supply a reduction from control problems in elections with a polynomial and binary positional scoring rule. We are thus provided with a novel correlation between hedonic games and elections, from which we deduce some of our complexity results. The main complexity results are summarized in Table 1.

We then move our attention to majority games, where agents form a graph and approve sets in which they are connected to the majority of other vertices. We classify the distribution of coalitions comprising perfect partitions and the social welfare incurred by them. We also explore this distribution in dynamic settings, where edges in the graph randomly and independently appear, via Erdős-Rényi graphs (Erdős and Rényi 1959). Regarding the standard measure of disutility caused by selfish behavior, Price of Anarchy (Koutsoupias and Papadimitriou 1999; Elkind, Fanelli, and Flammini 2016), we supply upper and lower bounds on the worst-case ratio between the (expected) social welfare of a partition maximizing it and that of a perfect one.

Related Work

The study of hedonic games was initiated by Dreze and Greenberg (1980), and later expanded to the study of characterizing various solutions concepts (concerning stability, perfection, and optimality), as well as many classes of hedonic games (See (Aziz, Savani, and Moulin 2016; Woeginger 2013) for surveys on the topic). As mentioned earlier, hedonic games with dichotomous preferences are one natural, succinctly representable class proposed by Aziz et al. (2016), who treat the exponential space requirement by representing agents' preferences by propositional formulas, and thus term such games as Boolean hedonic games. Peters (2016) thoroughly studied the complexity of finding optimal and stable partitions for various classes of DHGs.

Prior work on identifying solution concepts in hedonic games under uncertainty assumes the existence of

user preference data over some coalition, which is then used to construct probably approximately stable outcomes (termed as PAC stability). (Sliwinski and Zick 2017; Igarashi, Sliwinski, and Zick 2019) aim at learning players' preferences from data and obtain a PAC approximation of the original hedonic game, and then finding a partition that PAC-stabilizes the approximate hedonic game. In contrast, Jha and Zick (2020) focus on directly learning a variety of economic solution concepts from data. We remark that this line of research relies on (Balcan, Procaccia, and Zick 2015), who focus on transferable utility (TU) cooperative games.

Built upon the notorious network reliability problem (Provan and Ball 1983), the reliability extension of TU cooperative games are proposed by (Bachrach et al. 2011; Bachrach, Kash, and Shah 2012), encapsulating the effects of independent agent failures. They show how to approximate the Shapley value in such games using sampling, and how to accurately compute the core in games with few agent types. They further show that applying the reliability extension may stabilize the game. (Bachrach and Shah 2013) extend this model to weighted voting games, and propose algorithms for computing the value of a coalition, finding stable payoff allocations, and estimating the power of agents. These studies are strictly contrasted to our own work, since hedonic games are with non-transferable utility.

Focusing on symmetric friend-oriented hedonic games, where players have strong favour towards their friends, Igarashi et al. (2019) explore how stability can be maintained even after any set of at most k players leave their coalitions. They establish a robustness concept, which is close to the notion of fault tolerance in the theory of distributed systems (Fedoruk and Deters 2002). Although their work is most similar to ours, our work differs considerably. First, to the best of our knowledge, no attempt has been ever made to connect hedonic games and probabilistic inference. We thus aim at making the first step towards bridging this gap. Further, we also consider uncertainty of agents' friendships, instead of just their participation. Finally, we extend their model to a general probabilistic setting, whereas

they concern participation in a uniform and deterministic sense. This extension is inspired by (Imber and Kimelfeld 2021; Wojtas and Faliszewski 2012), who concern the problem of computing the probability of winning in an election where voter attendance is uncertain. In both papers, this task reduces to counting variants of election control problems (Bartholdi et al. 1992), whose goal is ensuring that a preferred candidate is the winner by controlling the set of either voters or candidates. Particularly, we observe that similar control schemes arise in our own framework.

Preliminaries

For an integer $n > 0$, let $[n] := \{1, \dots, n\}$. A hedonic game (HG) $\mathcal{G} = \langle N, (\succeq_i)_{i \in N} \rangle$ is given by a finite set $N = [n]$ of n agents, with a complete and transitive preference relation \succeq_i over $\mathcal{N}_i = \{C \subseteq N : i \in C\}$ for each agent $i \in N$. For each agent $i \in N$, we let \succ_i and \sim_i be the strict and indifference parts of \succeq_i (resp.). The outcome of a hedonic game \mathcal{G} is a partition π of N into disjoint coalitions. Let $\pi(i)$ be the coalition $C \in \pi$ such that $i \in C$. In a dichotomous hedonic game (DHG), agents only approve or disapprove coalitions, i.e., for each agent $i \in N$ there exists a utility function $v_i : \mathcal{N}_i \rightarrow \{0, 1\}$ such that $\pi(i) \succeq_i \pi'(i)$ iff $v_i(\pi(i)) \geq v_i(\pi'(i))$. We use the convention that $v_i(C) = 0$ whenever $i \notin C$ for some $C \subseteq N$.

Inspired by recent work on voting games (Imber and Kimelfeld 2021), we study a new variant of DHGs, referred to as Random DHGs (RDHGs), where each agent has a fixed and independent participation probability for being randomly drawn from N and, hence, of having the opportunity for forming coalitions with the other participants. Sometimes, the participation probability of an agent is unknown, and we thus briefly discuss in Appendix A how it can be estimated (Omitted and full proofs are available in (Cohen and Agmon 2023)). Formally, let $(p_i)_{i \in N} \in [0, 1]^n$ be the probabilities. Let $I \subseteq [n]$ be a random variable, where each $i \in N$ is in I with probability p_i and different indices are independent. The probability of I being a subset $U \subseteq [n]$ is $\Pr[I = U] = \prod_{i \in U} p_i \prod_{i \in [n] \setminus U} (1 - p_i)$. The RDHG induced by a DHG \mathcal{G} is thus $\mathcal{G}' = \langle N', (\succeq_i)_{i \in I} \rangle$, where the random set of players that participate in the game is $N' = I$. Agents won't change their true preferences over coalitions depending on the realization of the participants. Thus, agents (sincerely) value a coalition as if all its members were participating even if it is possible that they are not. That is, letting $v_i : \mathcal{N}_i \rightarrow \{0, 1\}$ be agent i 's utility function in \mathcal{G} , her utility function $v'_i : \mathcal{N}_i \rightarrow \{0, 1\}$ in \mathcal{G}' satisfies $v'_i(C) = v_i(C \cap I)$ for $C \in \mathcal{N}_i$. We note that v'_i may differ from v_i as there may be a coalition approved by agent i in \mathcal{G} that is disapproved in the induced RDHG due to the participants' realizations.

Given a partition or a coalition under our randomized setting, we redefine herein different measures of optimality and stability. Specifically, we largely follow

the solution concepts investigated by (Peters 2016). Given a partition π , a set \mathcal{C} of coalitions and a subset of players $S \subseteq N$, the social welfare of a coalition $C \in \mathcal{C}$ in the game restricted to the players S is $\mathcal{SW}_S(C) = \sum_{i \in C \cap S} v_i(C)$. $C \in \mathcal{C}$ is welfare-optimal w.r.t. \mathcal{C} ($\text{WO}_{\mathcal{C}}$) iff $C \in \arg \max_{\tilde{C} \in \mathcal{C}} \mathcal{SW}_S(\tilde{C})$. The social welfare of π is $\mathcal{SW}_S(\pi) = \sum_{C \in \pi} \mathcal{SW}_S(C)$. π is welfare-optimal (WO) iff $\pi \in \arg \max_{\tilde{\pi}} \mathcal{SW}_S(\tilde{\pi})$. A partition π is perfect (PF) if every agent is in an approved coalition in π . For stability, a partition π is core-stable if there is no non-empty coalition $C \subseteq N$ with $C \succ_i \pi(i)$ for all $i \in C$, and is strict-core-stable (SCS) if there is no non-empty coalition $C \subseteq N$ with $C \succeq_i \pi(i)$ for all $i \in C$ and $C \succ_i \pi(i)$ for some $i \in C$. In both, a group of agents may deviate. If we restrict our attention to the possibility of just a single agent deviating, we obtain the notion of Nash-stability (NS). That is, $\pi(i) \succeq_i \pi(j) \cup \{i\}$ for all i, j and π is individually rational (IR), i.e., $\pi(i) \succeq_i \{i\}$ $\forall i$.

For a RDHG \mathcal{G}' , sets \mathcal{C} and Π of polynomially many coalitions and partitions (resp.), i.e., $|\mathcal{C}|$ and $|\Pi|$ are $O(\text{poly}(n))$, a coalition $C \subseteq N$ and a partition π of N , we define the following events under \mathcal{G}' (while setting $S = I$ in \mathcal{SW}_S): $\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)$ denotes that $C \in \mathcal{C}$ is $\text{WO}_{\mathcal{C}}$; $\mathcal{WO}_{\mathcal{G}'}^{\Pi}(\pi)$ denotes that $\pi \in \Pi$ is WO_{Π} ; $\mathcal{WO}_{\mathcal{G}'}$ denotes that there is a WO partition; $\overline{\mathcal{WO}_{\mathcal{G}'}}^{\mathcal{C}}(C)$ denotes that C is not $\text{WO}_{\mathcal{C}}$; $\mathcal{NS}_{\mathcal{G}'}(\pi)$ denotes that π is NS; $\mathcal{NS}_{\mathcal{G}'}$ denotes that there is an NS partition; $\mathcal{PF}_{\mathcal{G}'}$ denotes that π is perfect; $\mathcal{PF}_{\mathcal{G}'}$ denotes that there is a perfect partition; $\mathcal{SCS}_{\mathcal{G}'}$ denotes that there is a strict-core-stable partition. When \mathcal{C} and Π are omitted, we refer to all possible coalitions and partitions (resp.). Note that the events' dependence on a specific game \mathcal{G}' quantifies the random set of players $I \subseteq [n]$ as defined for \mathcal{G}' .

Control Schemes. Throughout the paper, our control schemes are inspired by prior studies on winners in elections (Bartholdi et al. 1992; Hemaspaandra, Hemaspaandra, and Rothe 2007; Wojtas and Faliszewski 2012). Specifically, we focus on constructive control by adding players, where the goal is ensuring that a specified outcome satisfies a certain solution concept. Such a scheme can be employed by a central authority, which attempts to manipulate a DHG's outcome. Unlike prior studies on strategyproofness in hedonic games (Flammini et al. 2021), all players remain sincere in their preferences. Note that the following problems are not restricted to a specific class of DHGs.

Problem 1. We are given a DHG \mathcal{G} , a set of elements (either coalitions or partitions) \mathcal{C} , a set of players \mathcal{M} that already participate, a set \mathcal{Q} of non-participating players, an element $C \in \mathcal{C}$ and a solution concept β (e.g., welfare-optimality, perfection, Nash-stability). In our model, when a random set of players is drawn, there is no restriction on the number of players. Thus, in constructive control by adding an unbounded number of players (β -#CCAUP) the goal is counting the number of sets $\mathcal{Q}' \subseteq \mathcal{Q}$ such that C satisfies the solution concept

β w.r.t. the set of players $\mathcal{M} \cup \mathcal{Q}'$.

Problem 2. β - \exists CCAUP: given \mathcal{C} , \mathcal{M} , \mathcal{Q} and β , the goal is counting the number of all subsets $\mathcal{Q}' \subseteq \mathcal{Q}$ s.t. there exists some $C \in \mathcal{C}$ that satisfies β w.r.t. the set of players $\mathcal{M} \cup \mathcal{Q}'$.

Problem 3. (Decision Variants) Let β - $\#P$ be one of the counting problems in Problems 1–2. In the decision problem β - \mathcal{P} corresponding to β - $\#P$, instead of asking for a particular quantity w.r.t. to a solution concept β we ask if that quantity is greater than zero.

Problem 4. For each problem \mathcal{P} in Problems 1–3, in the problem \mathcal{P} - m we bound the number of selected players by a nonnegative integer $m \in \mathbb{N}$.

We recall that the class of counting variants of NP-problems is called $\#P$ and the class of functions computable in polynomial time is called FP. We also note the following:

Corollary 1. Let \mathcal{C} be a set of disjoint coalitions. Since different players are independent, the random social welfares of any pair of disjoint coalitions are independent, and thus different coalitions in \mathcal{C} are independent.

Reductions from CCAUP to RDHGs

Throughout our work, hardness results are derived from the following reductions from Problems 1–4 to RDHGs.

Lemma 1. There is a reduction from β - $\#$ CCAUP for coalitions to computing the probability that a coalition satisfies β . Thus, $\#P$ -completeness of β - $\#$ CCAUP implies $\#P$ -hardness of computing the mentioned probability.

Proof. Adapting the approach in Subsection 3.2 of (Imber and Kimelfeld 2021), let \mathcal{G} , \mathcal{C} , \mathcal{M} , \mathcal{Q} and $C \in \mathcal{C}$ be an instance of β - $\#$ CCAUP, where \mathcal{C} contains polynomially many disjoint coalitions. The players in \mathcal{M} participate with probability 1, whereas those in \mathcal{Q} participate with probability 1/2. Let $|\mathcal{Q}| := m$ and $\alpha(\mathcal{M}, \mathcal{Q})$ be the number of subsets $\mathcal{Q}' \subseteq \mathcal{Q}$ s.t. C satisfies β w.r.t. $\mathcal{M} \cup \mathcal{Q}'$. Corollary 1 yields that the probability that a coalition satisfies β equals to $2^{-m} \alpha(\mathcal{M}, \mathcal{Q})$, which concludes the proof. \square

Lemma 2. There is a reduction from β - \exists CCAUP for partitions to deciding the zeroness of the probability that there exists a partition satisfying β . Thus, NP-completeness of β - \exists CCAUP implies NP-hardness of the later.

Proof. Let \mathcal{G} , \mathcal{C} , \mathcal{M} , \mathcal{Q} be an instance of \exists CCAUP for partitions. The players in \mathcal{M} participate with probability 1 and those in \mathcal{Q} participate with probability 1/2. Clearly, the CCAUP instance admits a solution iff the probability that the constructed RDHG \mathcal{G}' admits a β partition is positive. \square

Thus, we hereafter focus on β - $\#$ CCAUP and β - \exists CCAUP for the sake of analyzing the complexity of probabilistic inference under RDHGs.

Algorithm 1: Computing $\mathcal{L}(n, \{q_m\}_{m \in [M]})$

```

1:  $\mathcal{L}(0, \{0\}_{m \in [M]}) = 0$ 
2:  $\mathcal{L}(0, \{q_m\}_{m \in [M]}) = 1$  ( $q_m \neq 0 \ \forall m \in [M]$ )
3: for  $j = 1$  to  $n$  do
4:   if  $\exists \tilde{m} \in [M]$  s.t.  $v_j(C_{\tilde{m}}) = 1$  then
5:      $\mathcal{L}(j, \{q_m\}_{m \in [M]}) = p_j \mathcal{L}(j-1, \{q_m\}_{\tilde{m} \neq m \in [M]} \cup \{q_{\tilde{m}} - 1\}) + (1 - p_j) \mathcal{L}(j-1, \{q_m\}_{m \in [M]})$ 
6:   else
7:      $\mathcal{L}(j, \{q_m\}_{m \in [M]}) = \mathcal{L}(j-1, \{q_m\}_{m \in [M]})$ 
8:   end if
9: end for
10: return  $\mathcal{L}(n, \{q_m\}_{m \in [M]})$ 

```

Hedonic Games with $\text{poly}(n)$ -Lists

Peters (2016) refers to a context where agents only approve polynomially many coalitions, and thus their preferences can be represented by merely listing all approved coalitions. In an even more restricted variant, the k -lists representation is considered, where each agent submits a list of at most a constant number of $k \in \mathbb{N}$ approved coalitions. In this section, we therefore explore the complexity of probabilistically inferring optimality in k -lists. We first investigate the complexity of computing the probabilities concerning welfare-optimal coalitions/partitions. Provided that it is $\#P$ -hard for k -lists with $k \geq 2$ (Theorem 3), we discuss their approximability. Finally, we prove that deciding the zeroness of the probability that there exists a SCS partition is NPC even for 1-lists (Theorem 6). Note that a hardness result for 1-lists also applies to k -lists with $k \geq 2$ (or even $\text{poly}(n)$ -lists).

Probability of Welfare-Optimality

For 1-lists, we prove that the probability that a coalition is welfare-optimal w.r.t. either a set \mathcal{C} or a set Π of polynomially many coalitions and partitions (resp.), and $\#$ CCAUP for welfare-optimal coalitions are poly-time computable.

Theorem 1. Computing $\Pr[\mathcal{WO}_{\mathcal{G}}^{\mathcal{C}}(C)]$, $\Pr[\mathcal{WO}_{\mathcal{G}}^{\Pi}(\pi)]$, as well as solving $\#$ CCAUP(- m) (Problems 1, 4) for welfare-optimal coalitions, can all be done in poly-time for 1-lists.

Proof. (Sketch) Note that $\Pr[\mathcal{WO}_{\mathcal{G}}^{\mathcal{C}}(C)] = \sum_{j=0}^n \Pr[\mathcal{WO}_{\mathcal{G}}^{\mathcal{C}}(C) \cap \mathcal{SW}_I(C) = j] = \sum_{j=0}^n \Pr[\mathcal{SW}_I(C) = j \cap [\cap_{C \neq C' \in \mathcal{C}} \mathcal{SW}_I(C') \leq j]]$. If \mathcal{C} contains only disjoint coalitions (i.e., $C \cap C' = \emptyset$ for any pair $C \neq C'$ in \mathcal{C}), then Corollary 1 can be invoked. Hence, this assumption drastically simplifies the proof, as discussed in Appendix B. Yet, we herein consider the more general case where this assumption is not necessarily satisfied.

Thus, we depict how $\Pr[\mathcal{SW}_I(C) = q \cap [\cap_{C \neq C' \in \mathcal{C}} \mathcal{SW}_I(C') \leq q]]$ can be computed in poly-time for any integer $0 \leq q \leq n$ and coalition $C \in \mathcal{C}$. We

denote $\mathcal{C} := \{C_1, \dots, C_M\}$, where $M = O(\text{poly}(n))$ by our assumption. For integers $j \in [n]$, $0 \leq q_m \leq n$ ($m \in [M]$), let $\mathcal{L}(j, \{q_m\}_{m \in [M]}) = \Pr[\cap_{m \in [M]} \sum_{i \in I \cap [j]} v_i(C_m) = q_m]$ (Recall that $v_i(C_m) = 0$ if $i \notin C_m$). In Appendix C, we observe that it can be computed in poly-time via the dynamic program in Algorithm 1. Note that $\mathcal{L}(n, \{q_m\}_{m \in [M]})$ is the desired probability. By summing these values, we can thus readily compute $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)]$ as described above. Computing $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\Pi}(\pi)]$ is via similar arguments, and thus deferred to Appendix D. Solving either $\# \text{CCAUP-}m$ or $\# \text{CCAUP}$ for welfare-optimal coalitions is by slightly modifying the proof of Theorem 6 in (Wojtas and Faliszewski 2012), which is thereby illustrated in Appendix E. \square

In contrast to Theorem 1, computing $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)]$ is generally intractable for k -lists with $k \geq 2$. Initially, we supply a generic reduction from winner elections to poly(n)-lists (See Appendix F.1 for a brief on voting games).

Theorem 2. There exists a poly-time reduction from each control problem in elections with a polynomial and binary positional scoring rule to its parallel control problem in poly(n)-lists under welfare-optimal coalitions.

Proof. (Sketch) The proof is for control by adding players only. When deleting players, the proof similarly follows via the same reduction. Let r be a scoring rule, \mathcal{C} be a set of candidates, $\mathcal{M} := \{M_{i_1}, \dots, M_{i_{h_1}}\}$ and $\mathcal{Q} := \{Q_{j_1}, \dots, Q_{j_{h_2}}\}$ be a voting profile of the registered and unregistered voters (resp.), and $c \in \mathcal{C}$ be a preferred candidate. If T_i is voter v_i 's ranking, then $s(T_i, c, r)$ is the score that the voter i contributes to a candidate c under r . We construct a DHG in poly(n)-lists form as follows. Let $\tilde{\mathcal{M}} := \mathcal{M} \cup \mathcal{C}$ and $\tilde{\mathcal{Q}} := \mathcal{Q}$ be the sets of participating and non-participating players (resp.). For each candidate $d \in \mathcal{C}$, we construct a coalition U_d containing d and all voters $v_{i_{\ell_1}}$ and $v_{j_{\ell_2}}$ ($\ell_1 \in [h_1], \ell_2 \in [h_2]$), which approve d . Formally, $U_d = \tilde{\mathcal{M}}_d \cup \tilde{\mathcal{Q}}_d \cup \{d\}$ with $\tilde{\mathcal{M}}_d = \{v_{i_{\ell_1}} : s(M_{i_{\ell_1}}, c, r) = 1\}$ and $\tilde{\mathcal{Q}}_d = \{v_{j_{\ell_2}} : s(Q_{j_{\ell_2}}, c, r) = 1\}$. Each candidate d submits an empty list and each voter v_i in the constructed hedonic game approves all coalitions U_d for which $v_i \in U_d$. Since $v_i \in U_d$ if and only if v_i approved the candidate d , then each player in the constructed hedonic game approves poly(n) many coalitions as r is a polynomial scoring rule. Thus, the reduction is well-defined. Finally, let $\tilde{\mathcal{C}} = \{U_d\}_{d=1}^m$ be the set of all possible coalitions and consider U_c as the preferred coalition. In Appendix F.2, we prove that every subset of voters $\mathcal{Q}' \subseteq \mathcal{Q}$ s.t. c is a winner of $\mathcal{M} \circ \mathcal{Q}'$ under r corresponds one-to-one to a subset of players $\tilde{\mathcal{Q}}' \subseteq \tilde{\mathcal{Q}}$ s.t. U_c is welfare-optimal w.r.t. $\tilde{\mathcal{M}} \cup \tilde{\mathcal{Q}}'$. \square

Theorem 2 thus enables us to transfer prior results on

controlling elections to hedonic games. Theorem 3 illustrates its application to welfare-optimal coalitions. An alternate proof for k -lists ($k \geq 2$) appears in Appendix G.

Theorem 3. In both k -lists ($k \geq 2$) and poly(n)-lists, if \mathcal{C} contains polynomially many disjoint coalitions, then $\# \text{CCAUP}$ for welfare-optimal coalitions and computing $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)]$ are $\# \text{P-hard}$.

Proof. For a fixed $k \geq 2$, Theorem 2 provides a reduction to $\# \text{CCAUP}$ for welfare-optimal coalitions from constructive control by adding an unlimited number of voters ($\# \text{CCUAV}$) under k -approval, which is $\# \text{P-hard}$ due to Theorem 3.2 in (Imber and Kimelfeld 2021). For poly(n)-lists, Theorem 2 provides a reduction to $\# \text{CCAUP}$ for welfare-optimal coalitions from $\# \text{CCUAV}$ under approval voting, which is $\# \text{P-hard}$ due to Theorem 13 in (Wojtas and Faliszewski 2012). By Lemma 1, computing $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)]$ is $\# \text{P-hard}$. \square

Despite Theorem 3, it appears that verifying $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)]$'s zeroness can be done in poly-time:

Theorem 4. In k -lists ($k \geq 2$) and poly(n)-lists, then CCAUP for welfare-optimal coalitions (partitions) and deciding $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)] > 0$ ($\Pr[\mathcal{WO}_{\mathcal{G}'}^{\Pi}(\pi)] > 0$) are in FP.

Proof. Let $\mathcal{G}, \mathcal{C}, \mathcal{M}, \mathcal{Q}$ and $C \in \mathcal{C}$ be an instance of CCAUP. Let $\mathcal{Q}^* \subseteq \mathcal{Q}$ be the set of all players which approve the coalition C . Let $\mathcal{Q}' \subseteq \mathcal{Q}$ s.t. C is welfare-optimal w.r.t. $\mathcal{M} \cup \mathcal{Q}'$. Then, C is welfare-optimal w.r.t. $\mathcal{M} \cup \mathcal{Q}^*$, and thus verifying this property is sufficient for deciding whether $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)] > 0$. Indeed, for each $i \in \mathcal{Q}^* \setminus \mathcal{Q}'$, adding i to $\mathcal{M} \cup \mathcal{Q}'$ increases the social welfare of C by 1, and the social welfare of any other coalition increases by at most 1. Thus, C is welfare-optimal w.r.t. $\mathcal{M} \cup \mathcal{Q}''$, where $\mathcal{Q}'' := \mathcal{Q}' \cup \mathcal{Q}^*$. Since the players in $\mathcal{Q}'' \setminus \mathcal{Q}^*$ disapprove C and the social welfare of the other coalitions cannot increase, C remains welfare-optimal even after removing them from $\mathcal{M} \cup \mathcal{Q}''$. Hence, C is welfare-optimal w.r.t. $\mathcal{M} \cup \mathcal{Q}^*$. The proof for $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\Pi}(\pi)] > 0$ is by similar arguments. \square

Remark 1. If we set $\mathcal{C} = \{C' \subseteq N : \exists i \in I \text{ s.t. } v_i(C') = 1\}$ and note that $|\mathcal{C}| \leq kn$, the results for computing $\Pr[\mathcal{WO}_{\mathcal{G}'}^{\mathcal{C}}(C)]$ also apply to computing $\Pr[\mathcal{WO}_{\mathcal{G}'}(C)]$.

Though Theorems 1 and 4 provided positive results, when attending to welfare-optimal partitions among all possible ones we achieve the following negative result for 1-lists.

Theorem 5. In 1-lists, $\# \exists \text{CCAUP}$ for welfare-optimal partitions w.r.t. Π is $\# \text{P-}$ and $\# \text{W}[1]$ -hard. Further, deciding $\Pr[\mathcal{WO}_{\mathcal{G}'}] > 0$ and $\exists \text{CCAUP}$ for welfare-optimal partitions are NPC and W[1]-hard.

Proof. (Sketch) Adapting the reduction in Theorem 5 of (Peters 2016), we show a reduction to $\# \exists \text{CCAUP}$ for welfare-optimal partitions of partitions

from #INDEPENDENT-SET (#IS), known to be #P-complete in general (Provan and Ball 1983). Given a graph $G = (V, E)$ and a target size k , we choose $\mathcal{M} = E$ and $\mathcal{Q} = V$ as the participating and the non-participating players (resp.). Each edge player $e \in E$ submits an empty list: she does not approve any coalition. A vertex player $v \in V$ approves $A_v := \{v\} \cup \{e \in E : v \in e\}$, i.e., v approves being together with the edges incident to it. In Appendix H, we prove that each independent set of size $\geq k$ corresponds one-to-one to a subset of players $\mathcal{Q}' \subseteq \mathcal{Q}$ and a partition π such that $\mathcal{SW}_{N'}(\pi) \geq k$ w.r.t. $N' = \mathcal{M} \cup \mathcal{Q}'$, thus providing us with a poly-time reduction. Our construction also supplies a reduction from INDEPENDENT-SET to CCAUP. By Lemma 2, deciding $Pr[\mathcal{WO}_{\mathcal{G}'}^C] > 0$ is #P-hard. Since #IS and IS are #W[1]- and W[1]-hard (Downey and Fellows 1995; Curticapean et al. 2019), respectively, the reduction also provides us with W[1]-hardness results where the parameter is the number of approving agents. \square

Approximate Probability of Welfare-Suboptimality

Similar to (Imber and Kimelfeld 2021), an additive Fully Polynomial-time Randomized Approximation Scheme (FPRAS) (Karger 1995) for $Pr[\mathcal{WO}_{\mathcal{G}'}^C(C)]$ can be obtained by a simple Monte Carlo estimation (by sampling and taking the ratio of the times in which C is welfare-optimal), whenever we can test in polynomial time whether C is welfare-optimal for a sample. Yet, a multiplicative FPRAS provides a stronger guarantee since it allows for approximating divisions of probabilities, which is required for estimating conditional probabilities. Thus, in Appendix I we depict the modifications of (Imber and Kimelfeld 2021)’s FPRAS for the probability of losing in an election so as to devise a FPRAS for the probability $Pr[\mathcal{WO}_{\mathcal{G}'}^C(C)]$.

Remark 2. (Disapproval k -Lists) Similar to k -lists, agents could only disapprove at most k coalitions. Thus, we note that disapproval k -lists also satisfy the above theorems.

Probability of Strict-Core-Stability

For strict-core-stability, we provide a negative result.

Theorem 6. For strict-core-stability, # \exists CCAUP is #P-hard in 1-lists. Further, \exists CCAUP and deciding $Pr[\mathcal{SCS}_{\mathcal{G}'}] > 0$ are NP-complete for 1-lists.

Proof. The proof is similar to that of Theorem 5, except that: (1) the reduction is from KERNEL (Szwarcfiter and Chaty 1994), the problem of counting kernels of a digraph (an independent set reachable from every outside node by an edge); and (2) each arc agent $e = (u, v)$ approves A_u . Due to space constraints, we omit the details. \square

Anonymous Hedonic Games

In an anonymous hedonic game (AHG) (Ballester 2004), agents’ preferences \succeq_i are determined by an underlying ordering \triangleright_i over the possible coalition sizes $[n]$, with $S \succeq_i T$ iff $|S| \triangleright_i |T|$. The following theorem analyzes the complexity of # \exists CCAUP and \exists CCAUP for Nash-stable, strict-core-stable and perfect partitions, as well as the verifying the zeroness of $Pr[\mathcal{NS}_{\mathcal{G}'}]$, $Pr[\mathcal{SCS}_{\mathcal{G}'}]$ and $Pr[\mathcal{PF}_{\mathcal{G}'}]$.

Theorem 7. For Nash-stability, strict-core-stability and perfection, # \exists CCAUP is #P-hard in AHGs. Further, \exists CCAUP, deciding $Pr[\mathcal{NS}_{\mathcal{G}'}]$, $Pr[\mathcal{SCS}_{\mathcal{G}'}]$, $Pr[\mathcal{PF}_{\mathcal{G}'}] > 0$, are all NP-complete in AHGs.

Proof. (Sketch) In Appendix J, we show a reduction from X3C to NS- \exists CCAUP that adapts the reduction from X3C presented in Theorem 9 of (Peters 2016). By Lemma 2, deciding $Pr[\mathcal{NS}_{\mathcal{G}'}] > 0$ is #P-hard. The proof for perfect- \exists CCAUP similarly adapts the reduction presented in Theorem 18 of (Woeginger 2013), and is thus deferred to Appendix J along with the proof for deciding $Pr[\mathcal{PF}_{\mathcal{G}'}] > 0$. The same reduction can be used for SCS. \square

Boehmer and Elkind (2020) show that finding a Nash-stable partition can be done in poly-time if each agent approves at most one coalition size. Despite the negative result in Theorem 7, we prove that such a condition further enables us to compute both $Pr[\mathcal{NS}_{\mathcal{G}'}(\pi)]$ and $Pr[\mathcal{PF}_{\mathcal{G}'}(\pi)]$ in polynomial-time for AHGs.

Theorem 8. If each agent $i \in N$ solely approves coalitions in $\mathcal{N}_i^+ = \{C \in \mathcal{N}_i : |C| = s_i\}$ for some $s_i \in \mathbb{N}$, then $Pr[\mathcal{NS}_{\mathcal{G}'}(\pi)]$ and $Pr[\mathcal{PF}_{\mathcal{G}'}(\pi)]$ are poly-time computable.

Proof. (Sketch) Note that $Pr[\mathcal{PF}_{\mathcal{G}'}(\pi)] = Pr[\cap_{i \in N} |\pi(i) \cap I| = s_i]$. If $|\pi(i)| < s_i$ for some $i \in N$, then $Pr[\mathcal{PF}_{\mathcal{G}'}(\pi)] = 0$. Thus, we hereafter assume that $|\pi(i)| \geq s_i$. Let $\mathbb{1}_i$ be the indicator for the event that agent i participates in the game. For integers $t \leq n$ and $0 \leq q_i \leq n$ ($i \in [n]$), let $\mathcal{L}(t, \{q_i\}_{i \in [n]}) = Pr[\cap_{i \in N} \sum_{j \in \pi(i) \cap I \cap [t]} \mathbb{1}_j = q_i]$. In Appendix K, we show that it can be computed in poly-time via the dynamic program in Algorithm 2, and so does the probability that π is perfect. For computing $Pr[\mathcal{NS}_{\mathcal{G}'}(\pi)]$, if there is no agent i with $s_i = 1$, then the grand coalition is NS. Thus, we hereafter assume that there exists at least one such agent. For each $j \in [n]$, let $N_j = \{i \in N : s_i = j\}$, and let $\ell = \max\{i | N_j = \emptyset \forall j \in [i]\}$. Boehmer and Elkind (2020) prove that for each $j \in [\ell]$ all agents in N_j need to be in coalitions of size j in every NS outcome. Noting that $Pr[\mathcal{NS}_{\mathcal{G}'}(\pi)] = Pr[\cap_{j \in [\ell]} \cap_{i \in N_j} |\pi(i) \cap I| = j]$, the proof thus follows from arguments similar to the previous one. \square

Algorithm 2: Computing $Pr[\mathcal{PF}_{G'}(\pi)]$

```

1:  $\mathcal{L}(0, \{0\}_{i \in [n]}) = 0, \mathcal{L}(0, \{q_i\}) = 1$  ( $q_i \neq 0 \forall i \in [n]$ )
2: for  $t = 0$  to  $n$  do
3:   if  $\exists i \in [n]$  s.t.  $\pi(t) = \pi(i)$  then
4:      $\mathcal{L}(t, \{q_i\}_{i \in [n]}) = p_t \mathcal{L}(t-1, \{q_i - 1\}_{i \in \pi(t)} \cup$ 
        $\{q_j\}_{j \in [n] \setminus \pi(t)}) + (1 - p_t) \mathcal{L}(t-1, \{q_i\}_{i \in [n]})$ 
5:   else
6:      $\mathcal{L}(t, \{q_i\}_{i \in [n]}) = \mathcal{L}(t-1, \{q_i\}_{i \in [n]})$ 
7:   end if
8: end for
9: return  $Pr[\mathcal{PF}_{G'}(\pi)] = \mathcal{L}(n, \{s_i\}_{i \in [n]})$ 

```

Candidate Intervals

Assuming the agent set can be placed in the natural ordering, each agent i only approves candidate intervals (CIs) $[a, b]$ of agents (with $i \in [a, b]$). Such a restriction was termed by Elkind and Lackner (2015), and applied to DHGs by (Peters 2016). Opposed to the negative result for k -lists with $k \geq 2$ (Theorems 3 and 5), we prove that:

Theorem 9. $Pr[\mathcal{WO}_{G'}^C(C)], Pr[\mathcal{WO}_{G'}^\Pi(\pi)]$ and $Pr[\mathcal{PF}_{G'}(\pi)]$ are computable in poly-time for CIs.

Proof. For an integer $0 \leq j \leq n$, let $\mathcal{SW}^*(j)$ and $\mathcal{SW}_I(j)$ be the maximum social welfare and the maximum coalitional social welfare over all coalitions (resp.) obtainable in the subgame restricted to the random agent set $I \cap [j]$. Each agent approves all originally approved coalitions S such that $S \subseteq I \cap [j] \cup \{0\}$. Note that $\mathcal{SW}^*(0) = 0$. Let $\#[t, j]$ be the number of agents that approve the interval $[t, j]$ in the subgame. Similar to Theorem 10 in (Peters 2016), we infer that $\mathcal{SW}^*(j) = \max_{t \in [j]} \{\mathcal{SW}^*(j-1) + \#[t, j]\}$, which can thus be computed in polynomial time via dynamic programming. Noting that $Pr[\mathcal{WO}_{G'}^\Pi(\pi)] = Pr[\mathcal{SW}_I(\pi) = \mathcal{SW}^*(n)]$ and that $Pr[\mathcal{SW}_I(\pi) = q]$ can be computed via dynamic programming (by arguments similar to the proof of Theorem 1, which are thus deferred to Appendix L), we infer that $Pr[\mathcal{WO}_{G'}^\Pi(\pi)]$ is in FP for CIs. By substituting $\mathcal{SW}^*(j)$ with $\mathcal{SW}_I(j)$, we obtain that $Pr[\mathcal{WO}_{G'}^C(C)]$ is also in FP for CIs. Since a partition π is perfect iff $\mathcal{SW}^*(\pi) = n$, the proof for $Pr[\mathcal{PF}_{G'}(\pi)]$ readily follows. \square

Roommate Games

In this section, we consider a restriction of hedonic games where agents only approve coalitions of size at most 2 (Gärdenfors 1975). The bipartite case of roommate games (RGs) is referred to as marriage games (Gale and Shapley 1962). See (Gusfield and Irving 1989) for a survey on both types of games. Finding perfect partitions is easy by Theorem 7 in (Aziz, Brandt, and Harrenstein 2013). Moreover, a core stable matching can be computed efficiently for marriage games (Gale and Shapley 1962). However, we obtain negative results

for both concepts in non-dichotomous marriage games and other classes of hedonic games.

Theorem 10. In marriage games, Perfect/Nash-stable- \exists CCAUP(- m) and deciding $Pr[\mathcal{NS}_{G'}], Pr[\mathcal{PF}_{G'}] > 0$ are NPC. The same applies to roommates, Representation by Individually Rational Lists of Coalitions (RIRLC) (Ballester 2004), additively separable hedonic games (Olsen 2009), B-hedonic games and W-hedonic games (Aziz et al. (2012)).

Proof. (Sketch) In Appendix M, we show a reduction from MinMaxMatch, the problem of finding a maximal matching with size $\leq m$, which is known to be NP-complete even for subdivision graphs (Horton and Kilakos 1993). The proof constitutes an adaptation of Theorem 1 in (Aziz 2013). The same reduction can be used for perfection. By Corollary 1 in (Aziz 2013), we further deduce the result for the other classes of hedonic games. \square

Perfect Partitions in Majority Games

Thus far, we analyzed the probability that either coalitions or partitions satisfy a solution concept. In contrast, we herein explore their average performance in majority games (Peters 2016). This class can be seen as a dichotomization of fractional hedonic games (Aziz et al. 2019). Formally, let $G = (N, E)$ be an undirected graph, where each agent corresponds to a vertex and an edge between two agents depicts a (mutual) friendship between them. Let $G_I = (I, E^I)$ be the random subgraph of G induced by I . Letting $d_i^I(C)$ be agent i 's degree in the subgraph of G_I induced by a coalition C , agent i approves C if $d_i^I(C) \geq \frac{|C|}{2}$, i.e., if i is connected to at least $\frac{|C|}{2}$ of the vertices in C . First, we characterize the distribution of both the coalitions and social welfare of perfect partitions (Theorems 11-12), on which we elaborate in Appendices N–O. Then, we discuss the performance of perfect outcomes by providing upper and lower bounds on their Price of Anarchy (Koutsoupias and Papadimitriou 1999).

By Theorem 14 in (Peters 2016), without loss of generality, a perfect partition consists of edges and triangles. Thus, we let M_n^I and T_n^I be the random variables which represent the number of edges and triangles in G_I (resp.). Let \mathbb{T} be the set of all triplets (i, j, k) ($i < j < k$) that form a triangle in G . Accordingly, the following theorem fully characterizes the social welfare of a perfect partition π , as well as the coalitions comprising π , for various values of $(p_i)_{i \in N}$.

Theorem 11. For each $i \in N$ and $n \in \mathbb{N}$, let $p_i(n) = \frac{q_i(n)}{n}$ for some $q_i : \mathbb{N} \rightarrow \mathbb{R}$, $q^{\max}(n) = \max_{i \in N} q_i(n)$ and $q^{\min}(n) = \min_{i \in N} q_i(n)$. Given a perfect partition π , we infer: (1) A perfect partition comprises of singletons w.h.p. (with high probability): If $q^{\max}(n) \rightarrow 0$ as $n \rightarrow \infty$, then $T_n^I = 0$, $M_n^I = 0$ and $\mathcal{SW}_I(\pi) = 0$; (2) Triangles and edges reside in perfect partitions a.s. (almost surely): If $|\mathbb{T}| = 1$ and $\frac{n}{q^{\min}(n)} \rightarrow 1$ as $n \rightarrow \infty$, then

$T_n^I \geq 1$ a.s. Otherwise, if $|\mathbb{T}| \geq 2$ and $q^{\min}(n) \rightarrow \infty$ as $n \rightarrow \infty$, then $T_n^I \geq 1$ and $M_n^I \geq 1$ a.s.; (3) If $\frac{q^{\max}(n)}{q^{\min}(n)} \rightarrow 1$ as $n \rightarrow \infty$, then $\frac{q^{\min}(n)}{n} \leq \mathbb{E}[|I|] \leq q^{\max}(n)$, thus yielding that $|I| \geq 1$ (i.e., at least one agent remains) a.s.; (4) If $q_i(n) \equiv c_i$ for $c_i > 0 \forall i$, then $\mathbb{E}[\mathcal{SW}_I(\pi)] \leq c_{\max}^2$, where $q^{\max}(n) \equiv c_{\max}$.

Further, we can model agents' uncertainty about their mutual friendships. Formally, let $(p_{ij})_{i,j \in N} \in [0, 1]^{n \times n}$ with $p_{ij} = p_{ji}$ for every $i, j \in N$. Let $\mathcal{E} \subseteq N \times N$ be a random variable, where $(i, j) \in \mathcal{E}$ with probability p_{ij} and different pairs of indices are independent, thus yielding a Erdős-Rényi random graph $\tilde{G} = (N, \mathcal{E})$ (Erdős and Rényi 1959) whose set of edges is \mathcal{E} . The majority game on the resulting random graph satisfies 1-3 in Theorem 11 with minor adjustments (See Appendix O.1), yet gives rise to an additional property which extends property 4 (proved in Appendix O.2):

Theorem 12. Let $p_{ij}(n) = c/n$ for some constant $c > 0$. Let π be a perfect partition. Then, T_n^I converges in distribution to a Poisson random variable with parameter $c^3/6$, $\mathbb{E}[M_n^I] = (n-1)c/2$ and $\mathbb{E}[\mathcal{SW}(\pi)] \leq (n-1)c$.

Let \mathbb{P}_I be the set of all perfect partitions for a random set of players I and let π^* be an welfare-optimal partition. Inspired by the Price of Anarchy (Koutsoupias and Papadimitriou 1999), we put forth the Price of Perfection (PP) of a RDHG \mathcal{G}' , defined as the worst-case ratio between the social welfare of π^* and that of a perfect partition, i.e., $\text{PP}(\mathcal{G}') = \max_{\pi \in \mathbb{P}_I} \frac{\mathcal{SW}_I(\pi^*)}{\mathcal{SW}_I(\pi)}$. Similarly, we define the Expected Price of Perfection (EPP) by $\text{EPP}(\mathcal{G}') = \max_{\pi \in \mathbb{P}_I} \frac{\mathbb{E}[\mathcal{SW}_I(\pi^*)]}{\mathbb{E}[\mathcal{SW}_I(\pi)]}$. Using Theorem 11, we devise upper and lower bounds on both variants of the price of perfection, where Corollary 2 is clearly a direct outcome of (3)-(4) in Theorem 11.

Lemma 3. Under the assumptions of Theorem 11 and: (1) in Theorem 11, $\text{PP}(\mathcal{G}') = \text{EPP}(\mathcal{G}') = 0$ w.h.p.; (2) in Theorem 11, $\text{PP}(\mathcal{G}') \leq |I|/2$ a.s.; (3)-(4) in Theorem 11, $\text{EPP}(\mathcal{G}') \leq q^{\max}(n) = c_{\max}/2$.

Proof. For (1), the claim clearly stems since $M_n^I = 0$ w.h.p. For (2), π^* clearly satisfies $\mathcal{SW}_I(\pi^*) \leq |I|$. If we were to consider each connected component of G separately, we may assume that G is connected and does not consists of any isolated vertices. Hence, if there exists a perfect partition in G , then a perfect partition consisting of edges and triangles exists (Theorem 14 in (Peters 2016)). However, G_I might contain isolated vertices, even if G does not. Since $M_n^I \geq 1$ a.s., we infer that $\mathcal{SW}_I(\pi) \geq 2$ at the very least, thus yielding that $\text{PP}(\mathcal{G}') \leq |I|/2$. For (3), we observe that $\mathbb{E}[\mathcal{SW}_I(\pi^*)] \leq \mathbb{E}[|I|] \leq q^{\max}(n)$. Combined with the proof for (2), we conclude that $\text{EPP}(\mathcal{G}') \leq q^{\max}(n)$. \square

Corollary 2. Under the assumptions of (3)-(4) in Theorem 11, if $\mathcal{SW}(\pi^*) \geq 1$, then $\text{EPP}(\mathcal{G}') \geq 1/c_{\max}^2$. Alternately, if $\mathcal{SW}_I(\pi^*) = |I|$, we infer that $\text{EPP}(\mathcal{G}') \geq \frac{c_{\min}}{nc_{\max}^2}$.

Proof. The first lower bound is a direct outcome of (4) in Theorem 11. For the second part, from (3) in Theorem 11 we infer that $\frac{q^{\min}(n)}{n} \leq \mathbb{E}[\mathcal{SW}_I(\pi^*)] \leq q^{\max}(n)$. Combined with (4) in Theorem 11, we conclude the desired bounds. \square

Conclusions and Future Work

Our work contributes significantly to the study of hedonic games, as the first one to explore the complexity of probabilistically inferring solution concepts in uncertain domains. The main complexity results are summarized in Table 1. Our study opens the way for many future works, including the investigation of other classes of hedonic games and other solution concepts. Further, our probabilistic setting arises several intriguing questions, among those: For an outcome satisfying a solution concept β , what is the maximum number of players whose withdrawal from the game still preserves β in the outcome induced by the remaining players? Another direction is robustness (Igarashi et al. 2019): A probabilistic withdrawal of players upon an outcome satisfying a solution concept β (e.g., stability) should preserve β .

Acknowledgements

This research was funded in part by ISF grant 1563/22.

We consider a multi-level jury problem in which ex

References

- Aziz, H. 2013. Stable marriage and roommate problems with individual-based stability. In *Proceedings of the 2013 international conference on Autonomous agents and multi-agent systems*, 287–294.
- Aziz, H.; Brandl, F.; Brandt, F.; Harrenstein, P.; Olsen, M.; and Peters, D. 2019. Fractional hedonic games. *ACM Transactions on Economics and Computation (TEAC)*, 7(2): 1–29.
- Aziz, H.; Brandt, F.; and Harrenstein, P. 2013. Pareto optimality in coalition formation. *Games and Economic Behavior*, 82: 562–581.
- Aziz, H.; Harrenstein, P.; Lang, J.; and Wooldridge, M. 2016. Boolean hedonic games. In *Proceedings of the 15th International Conference on Principles of Knowledge Representation and Reasoning*, 166–175.
- Aziz, H.; Harrenstein, P.; and Pyrga, E. 2012. Individual-based stability in hedonic games depending on the best or worst players. In *Proceedings of the 11th International Conference on Autonomous Agents and Multiagent Systems-Volume 3*, 1311–1312.
- Aziz, H.; Savani, R.; and Moulin, H. 2016. *Hedonic Games*, 356–376. Cambridge University Press.
- Bachrach, Y.; Kash, I.; and Shah, N. 2012. Agent failures in totally balanced games and convex games. In *International Workshop on Internet and Network Economics*, 15–29. Springer.
- Bachrach, Y.; Meir, R.; Feldman, M.; and Tennenholtz, M. 2011. Solving cooperative reliability games. In *Proceedings of the Twenty-Seventh Conference on Uncertainty in Artificial Intelligence*, 27–34.
- Bachrach, Y.; and Shah, N. 2013. Reliability weighted voting games. In *International Symposium on Algorithmic Game Theory*, 38–49. Springer.
- Balcan, M. F.; Procaccia, A. D.; and Zick, Y. 2015. Learning cooperative games. In *Twenty-Fourth International Joint Conference on Artificial Intelligence*.
- Ballester, C. 2004. NP-completeness in hedonic games. *Games and Economic Behavior*, 49(1): 1–30.
- Bartholdi, I.; John, J.; Tovey, C. A.; and Trick, M. A. 1992. How hard is it to control an election? *Mathematical and Computer Modelling*, 16(8-9): 27–40.
- Boehmer, N.; and Elkind, E. 2020. Individual-based stability in hedonic diversity games. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, 1822–1829.
- Cohen, S.; and Agmon, N. 2023. Complexity of Probabilistic Inference in Random Dichotomous Hedonic Games – Supplementary Material. <https://u.cs.biu.ac.il/~agmon/CohenAAAI23-Sup.pdf>.
- Curticapean, R.; Dell, H.; Fomin, F.; Goldberg, L. A.; and Lapinskas, J. 2019. A fixed-parameter perspective on #BIS. *Algorithmica*, 81(10): 3844–3864.
- Downey, R. G.; and Fellows, M. R. 1995. Fixed-parameter tractability and completeness II: On completeness for W [1]. *Theoretical Computer Science*, 141(1-2): 109–131.
- Dreze, J. H.; and Greenberg, J. 1980. Hedonic coalitions: Optimality and stability. *Econometrica*, 987–1003.
- Elkind, E.; Fanelli, A.; and Flammini, M. 2016. Price of Pareto optimality in hedonic games. In *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence*, 475–481.
- Elkind, E.; and Lackner, M. 2015. Structure in dichotomous preferences. *Proceedings of the 24th International Conference on Artificial Intelligence, IJCAI’15*, 2019–2025.
- Erdős, P.; and Rényi, A. 1959. On Random Graphs. I. *Publicationes Mathematicae Debrecen*, 6: 290–297.
- Fedoruk, A.; and Deters, R. 2002. Improving fault-tolerance by replicating agents. In *Proceedings of the first international joint conference on Autonomous agents and multiagent systems: part 2*, 737–744.
- Flammini, M.; Kodric, B.; Monaco, G.; and Zhang, Q. 2021. Strategyproof mechanisms for additively separable and fractional hedonic games. *Journal of Artificial Intelligence Research*, 70: 1253–1279.
- Gale, D.; and Shapley, L. S. 1962. College Admissions and the Stability of Marriage. *The American Mathematical Monthly*, 69(1): 9–15.
- Gärdenfors, P. 1975. Match making: assignments based on bilateral preferences. *Behavioral Science*, 20(3): 166–173.
- Gusfield, D.; and Irving, R. W. 1989. *The stable marriage problem: structure and algorithms*. MIT press.
- Hemaspaandra, E.; Hemaspaandra, L. A.; and Rothe, J. 2007. Anyone but him: The complexity of precluding an alternative. *Artificial Intelligence*, 171(5-6): 255–285.
- Horton, J. D.; and Kilakos, K. 1993. Minimum edge dominating sets. *SIAM Journal on Discrete Mathematics*, 6(3): 375–387.
- Igarashi, A.; Ota, K.; Sakurai, Y.; and Yokoo, M. 2019. Robustness against Agent Failure in Hedonic Games. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI-19*, 364–370.
- Igarashi, A.; Sliwinski, J.; and Zick, Y. 2019. Forming probably stable communities with limited interactions. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, 2053–2060.
- Imber, A.; and Kimelfeld, B. 2021. Probabilistic Inference of Winners in Elections by Independent Random Voters. *Proceedings of the 20th International Conference on Autonomous Agents and Multi-Agent Systems*, 647–655.
- Jha, T.; and Zick, Y. 2020. A learning framework for distribution-based game-theoretic solution concepts. In

Proceedings of the 21st ACM Conference on Economics and Computation, 355–377.

Karger, D. R. 1995. A randomized fully polynomial time approximation scheme for the all terminal network reliability problem. In Proceedings of the twenty-seventh annual ACM symposium on Theory of computing, 11–17.

Koutsoupias, E.; and Papadimitriou, C. 1999. Worst-case equilibria. In Annual Symposium on Theoretical Aspects of Computer Science, 404–413. Springer.

Olsen, M. 2009. Nash stability in additively separable hedonic games and community structures. *Theory of Computing Systems*, 45(4): 917–925.

Peters, D. 2016. Complexity of hedonic games with dichotomous preferences. *Proceedings of the AAAI Conference on Artificial Intelligence*, 30(1).

Provan, J. S.; and Ball, M. O. 1983. The complexity of counting cuts and of computing the probability that a graph is connected. *SIAM Journal on Computing*, 12(4): 777–788.

Sliwinski, J.; and Zick, Y. 2017. Learning Hedonic Games. In *IJCAI*, 2730–2736.

Szwarcfiter, J. L.; and Chaty, G. 1994. Enumerating the kernels of a directed graph with no odd circuits. *Information Processing Letters*, 51(3): 149–153.

Vallée, T.; Bonnet, G.; Zanuttini, B.; and Bourdon, F. 2014. A study of sybil manipulations in hedonic games. In 13th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2014), 21–28.

Woeginger, G. J. 2013. Core Stability in Hedonic Coalition Formation. In *SOFSEM 2013: Theory and Practice of Computer Science*, 33–50. Berlin, Heidelberg: Springer.

Wojtas, K.; and Faliszewski, P. 2012. Possible winners in noisy elections. In *Twenty-Sixth AAAI Conference on Artificial Intelligence*.