

Improved Knowledge Modeling and Its Use for Signaling in Multi-Agent Planning with Partial Observability

Shashank Shekhar

Department of Computer Science
Ben Gurion University
shekhar@cs.bgu.ac.il

Ronen I. Brafman

Department of Computer Science Information and Software Systems Eng.
Ben Gurion University
brafman@cs.bgu.ac.il

Guy Shani

Ben Gurion University
shanigu@bgu.ac.il

Abstract

Collaborative Multi-Agent Planning (MAP) problems with uncertainty and partial observability are often modeled as Dec-POMDPs. Yet, in deterministic domains, Qualitative Dec-POMDPs can scale up to much larger problem sizes. The best current QDec solver (QDec-FP) reduces MAP problems to multiple single-agent problems. In this paper we describe a planner that uses richer information about agents' knowledge to improve upon QDec-FP. With this change, the planner not only scales up to larger problems with more objects, but it can also support signaling, where agents signal information to each other by changing the state of the world.

1 Introduction

In many real-world problems, agents collaborate to achieve joint goals. For example, disaster response teams typically consist of agents that have multiple tasks to perform, some of which require the cooperation of several agents. In such domains, agents may have partial information, where they can sense only their immediate surroundings. As agents are often located in different positions and may possess different sensing abilities, their runtime information states differ. While this can be overcome using communication, the communication infrastructure can be damaged, or communication may be costly and should be reasoned about explicitly.

In this setting, it is common to compute a policy for all agents using a central engine. This policy is executed by the agents in a decentralized manner, and agent communication is performed through explicit actions. Decentralized POMDPs (Dec-POMDPs) offer a rich model for capturing such multi-agent problems (Bernstein et al., 2002; Oliehoek and Amato, 2016), but Dec-POMDP solvers have difficulty scaling up. Qualitative Dec-POMDPs (QDecs) were introduced as an alternative model, replacing the probability distributions over possible states with qualitative sets of states (Brafman, Shani, and Zilberstein, 2013). QDecs can also be viewed as the multi-agent extension of the well known Contingent Planning model (Hoffmann and Brafman, 2005). At least for deterministic problems, where partial ob-

servability plays the key role, QDec algorithms scale much better than Dec-POMDP algorithms. This was demonstrated clearly by the IMAP solver (Bazin and Shani, 2018), which was later improved by the QDec-FP planner (Shekhar, Brafman, and Shani, 2019).

QDec-FP solves an abstract and simplified centralized problem which they call the team problem. It is the single-agent problem, where the single agent is the entire team. This problem can be solved by a single-agent contingent solver. However, its solution is not necessarily executable by the agents. It may require that agent φ_j act differently based on the value of some proposition p observed by agent φ_i , even if φ_j cannot observe p . To fix this, a second set of planning problems is solved, in which each agent attempts to generate a plan that emulates its behavior in the solution to the team problem (which may require it to perform additional sensing, for example). If this second stage of QDec-FP is successful, we get a correct, executable solution. If it is not, the planner attempts to generate a new team plan.

In this paper we describe a new algorithm, QDec-FPS, which uses enhanced reasoning about the knowledge of individual agents in the team problem. Reasoning about individual agents' knowledge during team plan execution has two important advantages: First, it will lead to the generation of more informed team plans that are easier to extend to true solutions, because the team planner adds an action only if the agent that executes it knows that its preconditions hold. Second, it allows us to model, within the team plan, the process of explicit and implicit communication, which we refer to here as signaling.

Signaling refers to cases where agent φ_i can sense p but agent φ_j cannot. φ_i communicates this information to φ_j by setting the value of some variable q that agent φ_j can sense, to be correlated with the value of p . For example, φ_j cannot sense whether a door is open, but it can sense whether the light is on, while agent φ_i can sense both. If φ_i can also turn the light on and off, it can signal the door state to φ_j by turning on the light if and only if the door is open.

Technically, signaling consists of the following steps: (1) Agent φ_i senses p . (2) Agent φ_i sets the value of q to the value of p . (3) Agent φ_j senses q . To be sound, this

behavior must be consistent between the two execution branches that follow the sensing of p : if p is true, we must ensure q is true. If p is false, we must ensure q is false. Note that, to support signaling, the planner must model the knowledge of each agent within the team plan. For otherwise, the planner has no reason to insert signals into the plan because signaling does not enhance the knowledge of the team. Moreover, some problems cannot be solved without signaling.

2 Background

The flat-space QDec model is defined as follows:

Definition 2.1. A qualitative decentralized partially observable Markov decision process (QDec) is a tuple $\mathcal{Q} = \langle I, S, b_0, \{A_i | i \in I\}, \delta, \{\Omega_i\}, O, G \rangle$ where

- I is a finite set of agents indexed $1, \dots, m$. We often refer to the i^{th} agent as φ_i .
- S is a finite set of states.
- $b_0 \subset S$ is the set of states initially possible.
- A_i is a finite set of actions available to agent φ_i , and $\vec{A} = \otimes_{i \in I} A_i$ is the set of joint actions, where $\vec{a} = a_1, \dots, a_m$ denotes a particular joint action.
- $\delta : S \times \vec{A} \rightarrow 2^S$ is a non-deterministic Markovian transition function. $\delta(s, \vec{a})$ denotes the set of states the can be reached when taking joint action \vec{a} in state s .
- Ω_i is a finite set of observations available to agent φ_i and $\vec{\Omega} = \otimes_{i \in I} \Omega_i$ is the set of joint observation, where $\vec{o} = o_1, \dots, o_m$ denotes a particular joint observation.
- $\omega : \vec{A} \times S \rightarrow 2^{\vec{\Omega}}$ is a non-deterministic observation function. $\omega(\vec{a}, s)$ denotes the set of possible joint observations \vec{o} given that joint action \vec{a} was taken and led to outcome state s . Here $s \in S$, $\vec{a} \in \vec{A}$, $\vec{o} \in \vec{\Omega}$.
- $G \subset S$ is a set of goal states.

We focus on deterministic outcomes and deterministic observations. In such cases a successful solution is acyclic, and there is no need to bound the number of steps. We assume a shared initial belief, like most Dec-POMDP models.

A factored QDec representation is specified as follows: $\langle I, P, \{A_i | i \in I\}, Pre, Eff, Obs, b_0, G \rangle$. I is a set of agents, P is a set of primitive propositions, \vec{A} is a vector of individual action sets, Pre , Obs , and Eff are the precondition, observation, and effects functions, b_0 is the initial state formula, and G is a set (conjunction) of goal propositions.

The state space, S , consists of all truth assignments to P . Each state can be viewed as a set of literals. Initial states are those that satisfy the initial state formula. Goal states are those that satisfy the goal conjunction. The transition function, δ , is defined using Pre and Eff as follows: The precondition function Pre maps each individual action $a_i \in A_i$ to its set of preconditions, i.e., a set of literals that must hold whenever agent φ_i executes a_i . Preconditions are local, i.e., defined over

a_i rather than \vec{a} , because each agent must ensure that the relevant preconditions hold prior to executing its part of the joint action. We extend Pre to be defined over joint actions $\{\vec{a} = \langle a_1, \dots, a_m \rangle : a_i \in A_i\}$ (where $m = |I|$): $Pre(\langle a_1, \dots, a_m \rangle) = \cup_i Pre(a_i)$.

Following Bazinin and Shani (2018), we employ a more structured definition of joint actions. We assume that single-agent actions are not executed concurrently unless specified explicitly using collaborative actions. Collaborative actions have the same form as single-agent actions, except that they have multiple agent parameters. Thus, an agent may have a single-agent move action, as well as participate in a collaborative, two-agent action, joint-lift, for lifting a table. Later, the plan can be made more compact in post-processing, e.g., using the technique of Crosby, Jonsson, and Rovatsos (2014). For a deeper discussion of the definition of joint actions, see, e.g., (Shekhar and Brafman, 2018, 2020).

For every joint action \vec{a} and agent φ_i , $Obs(\vec{a}, i) = \{p_1, \dots, p_k\}$, where p_1, \dots, p_k are the propositions whose value agent φ observes after the joint execution of \vec{a} . The observation is private, i.e., each agent may observe different aspects of the world. We assume that the observed value is correct and corresponds to the post-action variable value.

We also assume that actions either sense or change the state of the world. Sensing actions do not affect the world state – their effect list is empty. Non-sensing actions provide no information – their observation list is empty. This is not a limiting assumption, as every action can be separated into a non-observation and a sensing action by adding suitable propositions forcing the two to appear consecutively in every plan. Finally, we assume every agent has a *noop* action, with $Pre(noop) = Eff(noop) = Obs(noop) = \emptyset$.

Another important concept in multi-agent planning is the notion of public and private actions and propositions (Brafman and Domshlak, 2008). A proposition that appears in the action descriptions of multiple agents is called public. Also, all goal propositions are public. An action is a public action if its precondition or effect contain a public proposition. A proposition that appears only in the actions of agent φ_i is called private to φ_i . An action whose preconditions and effects contain only private propositions is a private action.

The local plan of an agent φ is represented by a policy tree τ_i . Each node is labeled with an action, and edges that follow a sensing action are labeled by an observation. The agent performs the action at the root and continues to the subtree labeled with the observation it obtained.

Given policy tree τ_i for agent φ_i , and an o_i , $\tau_{i o_i}$ denotes the child subtree of the root of τ_i that is reached via a branch labeled by o_i . Let $\vec{\tau} = \langle \tau_1, \tau_2, \dots, \tau_m \rangle$ be a vector of policy trees, one per agent, also called a joint policy. We denote the joint action at the root of $\vec{\tau}$ by $\vec{a}_{\vec{\tau}}$, and for an observation vector $\vec{o} = o_1, \dots, o_m$, containing each agent's observation, we define $\vec{\tau}_{\vec{o}} = \langle \tau_{1 o_1}, \dots, \tau_{m o_m} \rangle$.

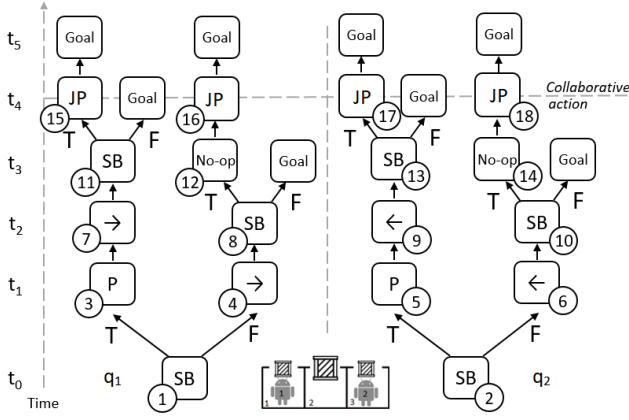


Figure 1: SB denotes sensing box, arrows denote motion, P denotes push, and JP denotes joint-push.

As actions may have preconditions, a joint policy tree is executable only if the preconditions of each action hold prior to its execution. To check this, we can maintain the set of states possible at each node during the execution of the joint policy, called the belief state. Initially, it consists of all possible initial states, denoted b_0 . After an agent with belief state b_i executes the first action of its current policy τ_i , a_{τ_i} , and observes o_i , its new belief state is $tr(b_i, o_i, a_{\tau_i}) = \{a_{\tau_i}(s) | s \in b_i, a_{\tau_i}(s) \models o_i\}$. Similarly the team belief state is initially b_0 . After executing the initial action of joint-policy τ given belief state b , it is updated to $tr(b, \vec{o}, \vec{a}_\tau) = \{\vec{a}_\tau(s) | s \in b, \vec{a}_\tau(s) \models \vec{o}\}$. Naturally, online, each agent may have less information than the combined team, as it does not know what other agents observed, and it cannot distinguish between all branches of the joint-policy. We denote by \vec{b} the set of agent-specific beliefs.

To follow policy $\vec{\tau}$, we first consider the action $\vec{a}_{\vec{\tau}}$ given the current belief state b . It must be the case that $b \models pre(a_{\vec{\tau}})$. In that case, we say that $\vec{a}_{\vec{\tau}}$ is executable in b . A joint policy $\vec{\tau}$ is executable in belief b if (1) $\vec{a}_{\vec{\tau}}$ is executable in \vec{b} ; (2) if a_{τ_i} is part of a collaborative action and j is another agent participating in that collaborative action, then a_{τ_j} contains j 's part of that action; (3) For every possible joint observation \vec{o} , $\vec{\tau}_{\vec{o}}$ is executable given $tr(b, \vec{o}, \vec{a}_{\vec{\tau}})$.

A joint policy is called a solution if it is executable in b_0 , and the belief state in all leaf nodes in the tree satisfies G . Note that unlike Dec-POMDPs, for QDecs there is no obvious notion of optimal policy, or optimization criterion, although one could strive to find trees with smaller depth, or trees that minimize the maximal branch cost.

Example 1. We have a 1D grid of size 3, with cells marked 1-3. The two agents φ_1 and φ_2 start in cells 1 and 3. Each cell may contain a box that needs to be pushed upwards. The left and right boxes are light and can be pushed by a single agent. The middle box is heavy and can be pushed by two agents acting together. Thus, $P = \{AgentAt_{i,pos}, BoxAt_{j,pos}, Heavy_j\}$

where $pos \in \{1, 2, 3\}$ is a grid position, $i \in \{\varphi_1, \varphi_2\}$, and $j \in \{1, 2, 3\}$ is a box index. In the initial state each box may be in or out of the grid — $b_0 = AgentAt_{\varphi_1,1} \wedge AgentAt_{\varphi_2,3} \wedge_{j=1,2,3} (BoxAt_{j,j} \vee \neg BoxAt_{j,j})$. There are therefore 8 possible initial states.

Agents can move left and right, push a light box up, or jointly push a heavy box up. There are no preconditions for moving left and right. The agent must be in the same cell as the box to push it. There also actions for sensing for boxes — $SenseBox_{\varphi,j}$, with precondition $Pre(SenseBox_{\varphi,j}) = AgentAt_{\varphi,j}$, no effects, and $Obs(SenseBox_{\varphi,j}) = BoxAt_{j,j}$. The goal is to move all boxes out of the grid, i.e., $\bigwedge_j \neg BoxAt_{j,j}$.

Figure 1 illustrates this domain and a possible solution.

In what follows projected sub-tree denotes a tree obtained from the original policy tree by removing some nodes. When a node n labeled by a non-sensing action is removed, the parent of n becomes the parent of the child of n . A node n labeled by a sensing action can only be removed if the two child subtrees of n are identical. In that case, the parent of n becomes the parent of one of the identical subtrees.

3 The QDec-FP Algorithm

QDec-FP (Shekhar, Brafman, and Shani, 2019) operates in three stages. (1) Generate a team solution. (2) Extend the projection of the team solution for each single agent. (3) Align the single agent plan trees. In the first step, the MA problem is transformed into a single-agent problem by treating all actions as if executed by the same meta-agent. This meta-agent can apply any agent's action and it is aware at each stage of the results of all sensing actions. We denote the resulting team plan as τ_{team} .

In the second step, from τ_{team} , QDec-FP generates for each agent φ_i a projection τ_{φ_i} . To obtain τ_{φ_i} we first remove from τ_{team} all non-sensing actions except those executed by agent φ_i . Next, we remove all private actions of all agents, leaving only nodes labeled by public actions. Furthermore, for each action a in τ_{φ_i} , we remove all its public preconditions, as they are guaranteed to be supplied by some public action in τ_{team} . Finally we remove redundant sensing actions — observations that do not influence how φ_i acts. These are sensing actions where both subtrees below the observation are identical, from the perspective of agent φ_i .

An example of a team solution, its projection, and the compacted projection is given in Figure 2. Figure 2A shows the team solution, containing both private and public actions of the agents. The team plan assumes shared knowledge. For example, in the left branch, only agent φ_1 senses for the existence of the heavy box, but then both agents jointly push it. Figure 2B shows the projected tree of φ_1 . Private actions of both agents were removed, but sensing actions of both agents remain, as no sensing action is redundant. Figure 2C shows the projected tree of φ_2 . The public push action executed

by φ_1 in the team plan was removed. And since φ_2 operates identically in both subtrees of the team plan, we remove the sensing action at the root of the team plan.

The projected tree is typically not executable by the agent. It contains sensing actions that are executed by other agents, and some actions require additional preconditions, due to the removal of private actions. Next, QDec-FP solves a single-agent contingent planning problem for each agent φ_i that requires φ_i to make τ_{φ_i} executable by adding sensing actions and private actions. QDec-FP does not add public actions, as they could interfere with the policies of the other agents by, e.g., consuming a precondition that is needed.

If some τ_{φ_i} is not solvable, QDec-FP backtracks and seeks a new team solution. If all τ_{φ_i} are solvable, we know that there is a solution, and all that remains is to align the different policies to ensure that agent actions are executed in the right order – as reflected by the team policy. This is done by inserting *noop* actions to postpone actions as need to be.

To implement QDec-FP, one needs a single-agent off-line contingent planner. QDec-FP uses CPOR (Kornitsky and Shani, 2016), which generates a policy tree by repeatedly calling the online contingent planner SDR (Shani and Brafman, 2011). SDR generates a single execution branch that corresponds to the true initial state. CPOR simply calls it multiple times with different "true" initial states.

While CPOR is the most scalable offline contingent planner currently, it has a major weakness: it does not backtrack. To overcome this, QDec-FP introduced a backtracking mechanism on top of CPOR. It modifies the planning problem with an additional, sound, constraint, that invalidates previous solutions. However, this limits the number of backtracks that are practically possible.

4 The QDec-FPS Planner

QDec-FPS has the same high-level structure as QDec-FP, but modifies the way the team problem is solved, and adds a mechanism for generating macros that enable signaling.

4.1 Agent Specific Knowledge

QDec-FP uses, through the CPOR planner, the SDR translation from contingent planning to classical planning. This translation maintains, for each proposition p , two propositions: Kp and $K\neg p$, denoting knowing that p is true or false, respectively. SDR then transforms each precondition p of an action to Kp . The translation ensures that Kp holds in a belief state only if p holds in all possible worlds. This is done by reasoning about all the specific possible worlds explicitly.

The state description in the classical translation contains the current state of the world given every possible initial state, in the form of $Kp|s$, where p is a proposition, and s is a possible state. It also contains actions

that allow deducing new knowledge facts, called merge and refutation actions (Palacios and Geffner, 2009). Thus, the agent can also obtain Kp if for every currently possible state s , $Kp|s$ holds.

Because QDec-FP solves the team problem as a single-agent planning problem, the planner treats all agents as part of a single centralized agent, and the knowledge of this agent reflects the combined knowledge of all agents. Thus, when one agent observes that p holds, other agents can use this knowledge, without observing the value of p themselves. This is inconsistent with the real plan tree, where each agent must independently ensure that p holds, before executing an action that requires p as precondition.

QDec-FPS changes SDR's translation to be agent aware. The Kp propositions used to capture combined knowledge are replaced by propositions of the form $K_\varphi p$, denoting that agent φ knows that p holds. The precondition p of an action executed by agent φ is replaced with $K_\varphi p$. The effect of a sensing action executed by φ is that only the sensing agent φ knows the value of p , i.e., either $K_\varphi p$ or $K_\varphi \neg p$. Similarly, merge and refutation actions now provide agent-specific inference. The effect of a non-sensing action is still known to all agents (as agents are aware of the plan itself).

This modification forces the underlying planner to insert sensing actions by different agents to ensure they have the knowledge required to perform their actions, whereas in QDec-FP, because all agents are treated as one, it was enough if some other agent performed this sensing action. If an agent has an action with a precondition p , the team plan will ensure that the agent first senses or learns the value of p . If the agent cannot sense or learn the value of p , such an action will not be part of a generated team plan.

This leads to the generation of team plans that better account for agent abilities. Like team plans in QDec-FP, these plans may not be extendable because they can require agents to act differently depending on information they do not have. Yet, they are much better informed, and are therefore less likely to lead to failure when agents extend the team plan. Hence, they often lead to fewer backtracks.

The new translation also allows us to solve problems that QDec-FP cannot solve. For example, imagine that Agent φ_1 can observe only p , and agent φ_2 can only observe q , but must execute an action with precondition p . QDec-FP's team plan will most likely let φ_1 sense p , and will make φ_2 execute the action afterwards. This team plan cannot be later extended by φ_2 to a valid local plan because it cannot sense p . QDec-FPS will avoid this team plan because it will know that φ_2 does not know p . Furthermore, as we discuss below, if φ_1 can set the value of q to be correlated with p , it will be able to exploit this implicit form of communication.

The new translation is more demanding and generates classical planning problems that are harder to solve. Because of this, it is able to detect unsolvable problems faster. However, in the special case of agents

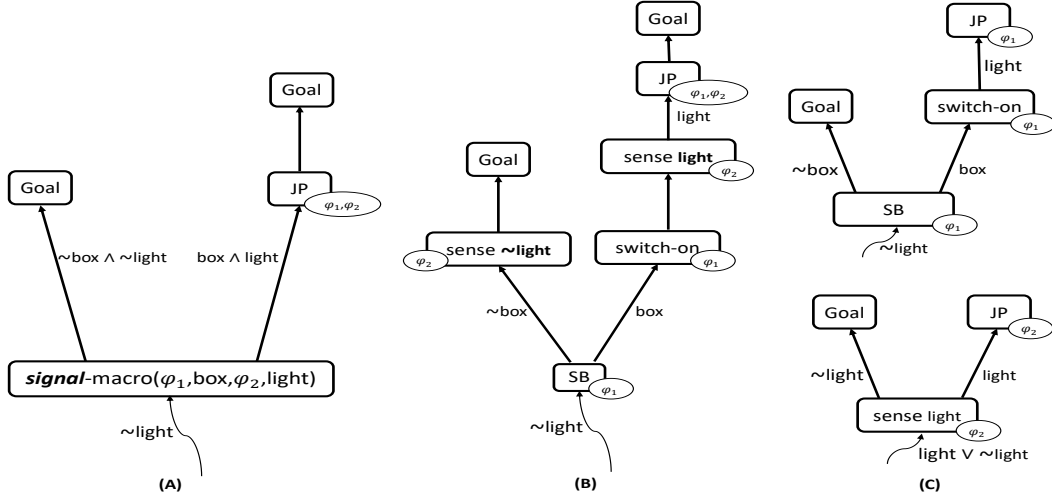


Figure 3: Signaling in the QDec-FPS planner. (A) A team plan with the signaling macro action. (B) The team plan following the macro expansion. (C) Projected trees for φ_1 (top) and φ_2 (bottom).

agents, φ_1 and φ_2 must jointly push a heavy box. φ_1 can sense the box, but φ_2 can only sense whether a light is on. φ_1 can signal to φ_2 about the box by turning on the light, which is originally off. Figure 3A shows a team plan with a macro. Figure 3B shows the team plan after macro expansion: φ_1 senses the box, then turns on the light, if needed, and then φ_2 senses the light. Finally, both jointly push the box. Figure 3C shows the projected single agent plan trees for φ_1 and φ_2 .

In practice, adding this macro on top of an online contingent solver that focuses on a single branch at a time, such as SDR, is not straightforward. To address this, we do the following: We add an action by φ_i that can be viewed as a commitment to ensure that $p \leftrightarrow q$ holds. This action is constrained to be followed immediately by the action of sensing p by φ_i . At this point, in the team plan, if agent φ_j needs to know the value of p , it can use the fact that $p \leftrightarrow q$ to deduce it from the value q . To ensure that it learns the value of p , we force the action of sensing q in both branches. As above, the team plan is post-processed to ensure that φ_i does indeed ensure the validity of $p \leftrightarrow q$ following the sensing action.

QDec-FPS Properties. When discussing soundness and completeness, we should first distinguish between the abstract approach used by QDec-FP/S, which reduces multi-agent contingent planning to single-agent contingent planning, and the practical implementation of the planner which uses a specific planner, CPOR.

First, consider the abstract formulation and assume a sound and complete underlying single-agent contingent planner. The soundness of QDec-FPS without signaling follows from that of QDec-FP: Every team solution of QDec-FPS is also a possible team-solution for QDec-FP and all the other steps are identical. This remains true in the implementation due to the soundness of CPOR.

With signaling, the abstract version of the macro is

a sensing action with effects. It remains sound if the action used by the signaling agent to ensure that $p \leftrightarrow q$ holds does not affect any other proposition. The implementation using CPOR is more complex, but it ensures that following the commitment to the signaling macro, the team solution implements a correct version of the macro, and hence it is sound, too.

In the case of completeness, for QDec-FPS, one cannot separate between an abstract and implemented case because the core contribution of QDec-FPS is at the level of the implementation of the classical encoding used by CPOR and SDR. At this level, QDec-FPS is incomplete for two reasons: First, CPOR is an incomplete solver: it lacks an internal backtracking mechanism. (For that reason, the implementation of QDec-FP is also incomplete, although the abstract formulation is complete.) Second, the planner’s ability to reason about the knowledge of separate agents, as implemented at present, is incomplete. That is, it can under-estimate the knowledge of an agent. Thus, it might not be able to deduce that the preconditions of some actions are known to the acting agent, even though they are.

5 Empirical Evaluation

We now examine the applicability and scalability of QDec-FPS by comparing it with QDec-FP on 3 multi-agent planning domains. Some problems in each domain were modified to require signaling. Both QDec-FP and QDec-FPS are implemented in C#, and were run on a Windows 10, 64 bit machine with i7 processor, 2.8GHz CPU, and 16Gb RAM. IMAP was not considered, as QDec-FP was shown to scale better than IMAP (Bazin and Shani, 2018).

We experiment with the following three domains: **Box-Pushing (BP)**: Boxes located in a grid-like structure. Each box is to be pushed to its destination location outside the column the box appears in (Brafman, Shani, and Zilberstein, 2013). A light box can be

Domain	Ins (#agt)	Objects	Max-width		Max-height		Time (sec)		BT	
			fp	fps	fp	fps	fp	fps	fp	fps
BP	B1 (3)	16	8	5	23	18	3.59	2.91	0	0
	B2 (4)	16	12	10	19	19	5.3	6.1	0	0
	B3 (9)	36	64	*	24	*	25.3	*	0	*
	B4 (3)	11	4	4	14	11	16.39	1.17	9	0
	B5 (3)	12	6	8	16	15	13.65	2.9	4	0
	B6 (3)	12	-	6	-	12	-	13.58	47+	4
	B7 (3)	12	8	8	18	19	158.89	3.87	41	0
	B8 (3)	12	8	8	17	21	111.6	4.05	26	0
	B9 (3)	13	16	14	21	19	121.42	5.6	19	0
	B10 (3)	16	16	15	26	29	155.83	9.69	33	0
	B11 (5)	20	*	24	*	32	*	75.21	*	1
	B12 (5)	20	*	24	*	37	*	365.9	*	6
	B13 (2)	10	na	2	na	6	na	1.06	na	0
	B14 (2)	12	na	2	na	2	na	1.20	na	1
	B15 (3)	12	na	4	na	12	na	2.49	na	0
	B16 (3)	13	na	4	na	7	na	5.5	na	1
	B17 (3)	13	na	4	na	6	na	8.9	na	2
	B18 (3)	14	na	6	na	16	na	4.04	na	0
	B19 (4)	15	na	8	na	14	na	17.6	na	2
TM	T1 (3)	10	8	7	20	16	2.68	2.78	0	0
	T2 (4)	12	16	13	21	17	7.84	8.26	0	0
	T3 (4)	15	12	16	34	26	8.74	12.39	0	0
	T4 (5)	14	19	22	22	24	20.5	43.58	0	0
	T5 (5)	16	12	14	32	25	9.7	11.2	0	0
	T6 (3)	8	2	2	8	8	11.01	0.61	7	0
	T7 (3)	10	8	7	20	16	34.5	41.8	8	8
	T8 (3)	10	8	8	24	22	37.87	22.45	9	4
	T9 (3)	10	-	-	-	-	140.32	29.73	31	6
	T10 (4)	12	16	16	23	22	6.68	152.18	0	14
	T11 (5)	16	16	16	30	35	274.5	56.15	27	3
	T12 (2)	10	na	2	na	9	na	1.7	na	0
	T13 (2)	12	na	4	na	17	na	6.81	na	0
	T14 (3)	12	2	2	16	11	17.59	2.32	9	0
	T15 (3)	13	na	4	na	14	na	7.80	na	0
	T16 (3)	14	na	4	na	19	na	11.68	na	1
Rovers	R1 (1)	12	2	2	8	8	3.8	3.8	0	0
	R2 (2)	14	2	2	9	8	5.4	5.3	0	0
	R3 (2)	13	4	4	15	15	9.3	8.9	0	0
	R4 (2)	17	12	12	34	23	35.8	38.6	0	0
	R5 (3)	18	12	12	28	32	51.5	54.1	0	0
	R6 (3)	21	30	30	36	29	136.6	111.9	0	0
	R7 (3)	23	98	81	50	45	567.9	233.2	0	0
	R8 (3)	13	2	2	7	6	57.9	5.9	3	0
	R9 (3)	14	4	4	12	12	113.8	10.2	4	0
	R10 (3)	14	4	4	11	11	314.5	89.3	11	3
	R11 (3)	19	-	12	-	12	325.1	41.1	6+	0
	R12 (4)	15	4	4	12	11	47.2	16.0	1	0
	R13 (4)	20	12	12	25	19	241.8	49.1	3	0
	R14 (2)	16	na	2	na	9	na	3.6	na	0
	R15 (2)	18	na	2	na	10	na	5.89	na	0
	R16 (2)	16	na	4	na	13	na	34.17	na	0
	R17 (2)	20	na	4	na	14	na	17.68	na	0
	R18 (3)	17	na	2	na	7	na	10.17	na	0
	R19 (3)	18	na	4	na	13	na	12.93	na	0
	R19 (3)	14	na	2	na	7	na	6.62	na	0
	R20 (3)	19	na	8	na	16	na	36.74	na	0
	R21 (4)	20	na	6	na	11	na	67.79	na	0
	R22 (4)	22	na	6	na	12	na	68.21	na	0

Table 1: Comparison of QDec-FP and QDec-FPS planners. Ins (#agt): instance number and number of acting agents. Object: the number of objects in each problem. BT: number of planner backtracks. *: time out. '-': planner could not solve or breaks down. na: problem not applicable to a solver. fp: QDec-FP. fps: QDec-FPS. The best approach, based on time only, is shown in bold.

pushed by a single agent co-located with the box while a heavy box requires two co-located agents. Agents can be non-homogeneous, i.e., different agents can observe and push different boxes.

Table-Mover (TM): TM consists of a number of tables and rooms, and agents that can move between connected rooms (Shekhar and Brafman, 2020). The tables' locations are uncertain initially, and agents must move them to their goal locations. Agents can be non-homogeneous in their sensing and manipulation abilities. All table manipulation actions are collaborative: 2move-table, 2lift-table, 2drop-table.

Rovers: Multiple rovers navigate a planet surface, finding samples and communicating them back to a lander

(Bazin and Shani, 2018). Two rovers must simultaneously collect the rock sample, while a single agent can sample soil as well as take images of certain objectives on its own. Coordination points include locations (waypoints) which are accessible to multiple rovers. Rovers communicate sampled soil/image/rock data to the lander that exists at a certain waypoint. A rover navigates between two waypoints and must be present at the corresponding to sample. Availability of data to sample at a waypoint is unknown to the rovers initially. In this modified domain, our schema requires two rovers working jointly to collect rock samples. After taking measurements, the rovers must broadcast them back to the lander. In this domain in general, a rover has fewer public actions (approx. 46% on average), but a relatively complex internal planning problem, including navigation, soil and rock sampling, and image capturing.

Table 1 compares QDec-FPS and QDec-FP based on policy quality (max-width, max-height), runtime (time), and the number of backtracks required. max-width and max-height refer to maximum number of branches and the maximum height of all individual solution trees obtained for the agents. The number of branches is also indicative of the number of sensing actions performed, as branching occurs following an observation. The planner backtracks when at least one of the single-agent problems, obtained by decomposing τ_{team} , is unsolved by CPOR. Within each domain in the table, dashed lines separate three problem classes: homogeneous agents, non-homogeneous agents, and non-homogeneous agents that require signaling. To handle signaling, QDec-FPS adds macros, which may have a significant overhead. Therefore, these macros were only added in problems that require signaling. The decision whether to add macros was done manually. In the future, we will automatically detect whether signaling is needed.

In BP, QDec-FPS scales much better than QDec-FP in problems with three to five agents. Increasing the number of objects had minor impact on QDec-FPS running time, as opposed to QDec-FP. For many problems, QDec-FPS needs to backtrack fewer times than QDec-FP, and as a result, it finds solutions faster. On the other hand, increasing the number of agents has an adverse effect on QDec-FPS. In fact, instance B3 in the BP domain with nine identical agents, was quite rapidly solved by QDec-FP, while QDec-FPS times out. This is because the new translation makes the team problem much harder to solve. Thus, QDec-FP finds a team plan quickly, and when agents are identical, it is usually very easy to fix the team problem by adding any needed sensing action. For that reason, in the case of homogenous agents, we see no backtracks in any of the problems. On the other hand, problems B11 and B12 were solved by the QDec-FPS planner but were unsolved by QDec-FP. This is most likely due to the high number of required backtracks. We can conclude that for simpler problems with identical agents, QDec-FP scales better with the number of agents, but for more

consider a multi-level jury problem in which experts We
consider a multi-level jury problem in which experts We
consider a multi-level jury problem in which experts

References

- Bazinin, S., and Shani, G. 2018. Iterative planning for deterministic QDec-POMDPs. In *GCAI 2018*, 15–28.
- Bernstein, D. S.; Givan, R.; Immerman, N.; and Zilberstein, S. 2002. The complexity of decentralized control of Markov decision processes. *Mathematics of Operations Research* 27:819–840.
- Brafman, R. I., and Domshlak, C. 2008. From one to many: Planning for loosely coupled multi-agent systems. In *ICAPS*, 28–35.
- Brafman, R. I.; Shani, G.; and Zilberstein, S. 2013. Qualitative planning under partial observability in multi-agent domains. In *Proceedings of the Twenty-Seventh AAAI Conference on Artificial Intelligence*, July 14–18, 2013, Bellevue, Washington, USA.
- Crosby, M.; Jonsson, A.; and Rovatsos, M. 2014. A single-agent approach to multiagent planning. In *ECAI 2014 - 21st European Conference on Artificial Intelligence*, 18–22 August 2014, Prague, Czech Republic - Including Prestigious Applications of Intelligent Systems (PAIS 2014), 237–242.
- Hoffmann, J., and Brafman, R. 2005. Contingent planning via heuristic forward search with implicit belief states. In *Proc. ICAPS*, volume 2005.
- Komarnitsky, R., and Shani, G. 2016. Computing contingent plans using online replanning. In *AAAI*, 3159–3165.
- Oliehoek, F. A., and Amato, C. 2016. *A Concise Introduction to Decentralized POMDPs*. Springer Briefs in Intelligent Systems. Springer.
- Palacios, H., and Geffner, H. 2009. Compiling uncertainty away in conformant planning problems with bounded width. *J. Artif. Intell. Res.* 35:623–675.
- Shani, G., and Brafman, R. I. 2011. Replanning in domains with partial information and sensing actions. In *IJCAI 2011, Proceedings of the 22nd International Joint Conference on Artificial Intelligence*, Barcelona, Catalonia, Spain, July 16–22, 2011, 2021–2026.
- Shekhar, S., and Brafman, R. I. 2018. Representing and planning with interacting actions and privacy. In *Proceedings of the Twenty-Eighth International Conference on Automated Planning and Scheduling*, *ICAPS 2018*, Delft, The Netherlands, June 24–29, 2018, 232–240.
- Shekhar, S., and Brafman, R. I. 2020. Representing and planning with interacting actions and privacy. *Artif. Intell.* 278.
- Shekhar, S.; Brafman, R. I.; and Shani, G. 2019. A factored approach to deterministic contingent multi-agent planning. In *Proceedings of the Twenty-Ninth International Conference on Automated Planning and Scheduling*, *ICAPS 2018*, Berkeley, CA, USA, July 11–15, 2019, 419–427.